

Sycophancy or Empathy ?

DeepReflect – An LLM-based system designed to analyze and generate responses to personal queries

Tara K. Jain

Supervisor: Dr. Diana Maynard

*A report submitted in fulfilment of the requirements
for the degree of MSc in Artificial Intelligence*

in the

Department of Computer Science

September 20, 2025

Declaration

All sentences or passages quoted in this report from other people's work have been specifically acknowledged by clear cross-referencing to author, work and page(s). Any illustrations that are not the work of the author of this report have been used with the explicit permission of the originator and are specifically acknowledged. I understand that failure to do this amounts to plagiarism and will be considered grounds for failure in this project and the degree examination as a whole.

Name:

Signature:

Date:

Abstract

This project explores how the same response can be interpreted as exhibiting two behaviors with divergent connotations in LLM-generated responses. In particular, the project demonstrates how the same generated response may be perceived as “sycophantic” or “empathic” depending on the psychosocial lens through which it is evaluated. We find that the issue is not isolated to LLM-generated responses and that human-authored ones are susceptible to dual interpretations too. Experiments across four language models (GPT-oss, GPT-4o, Claude, and Llama) demonstrate significant differences in the traits of LLM-generated responses compared to human-authored ones to the same personal query. The project uses posts and comments from popular subreddits (r/aitah and r/anxiety) to provide empirical evidence. The investigation indicates a strong association between certain traits of LLM responses which may be perceived as instances of “sycophancy” or “empathy”. The project operationalizes sycophantic behaviour with Goffman’s Theory of Face and empathic behaviour with Carl Rogers’ Person Centered Therapy framework. DeepReflect is the comparative system developed in this project for analyzing human and AI-generated responses to personal queries across multiple paradigms of psychosocial behavior. The framework is designed to be extensible, enabling it to incorporate additional psychosocial paradigms such as the Rokeach Value Survey (RVS) as additional instruments for analysis of the traits exhibited by responses. Additionally, it implements and investigates the effects of a control mechanism by inserting expert human transcripts to elicit improvements in the quality of responses with chain-of-thought (CoT) reasoning.

Contents

1	Introduction	1
1.1	Aims and Objectives	3
1.1.1	Research Questions	3
1.2	Key Contributions	4
1.3	Overview of the Report	4
2	Literature Survey	6
2.1	Psychosocial paradigms: Theoretical Foundations	7
2.1.1	Rogerian Person Centered Therapy (PCT)	7
2.1.2	Goffman’s theory of face (ToF)	8
2.1.3	Rokeach Value Survey (RVS)	8
2.2	Contemporary NLP research	9
2.2.1	LLM-as-judge paradigm	10
2.3	Gaps in the Literature and Open Challenges	10
3	DeepReflect: System Design	12
3.1	Dataset	12
3.1.1	Subreddit Selection	12
3.2	DeepReflect	13
3.2.1	System Design	13
4	Methods	18
4.1	Data preprocessing	18
4.2	Procedures	18
4.3	Experiments	19
4.3.1	Experiment 1	19
4.3.2	Experiment 2	20
4.3.3	Statistical Methods	20
4.4	CoT Generation Evaluation	20
4.5	Construct Validity and Evaluation Metrics	20

5	Results and Analysis	22
5.1	LLM-as-judge Annotation Reliability	22
5.2	"Sycophancy" or "Empathy" ?	23
5.2.1	Dual Coding of responses as "Empathic" and "Sycophantic"	23
5.2.2	Trends within human authored responses	24
5.2.3	Strength of association varies by model	26
5.3	Social Verdicts in r/aitah	27
5.3.1	Generations with CoT	28
5.3.2	Error analysis	29
5.3.3	Extensibility of the framework	29
6	Conclusion	31
6.1	Limitations	31
6.2	Summary and Conclusion	32
	Appendices	37
A	Appendix A	38
A.1	Annotation Prompts and Traits	38
A.2	Response Generation Prompts	38
A.2.1	Basic Response Generation Template	38
B	Analysis Reports	40

List of Figures

3.1	Pipeline architecture for DeepReflect.	14
3.2	CoT Generation with personal queries embedded in reasoning dialogs retrieved from expert human transcripts. In this case, the dialogs are from Carl Rogers’ sessions with Gloria (patient) [21]. This dialog was selected because it reflects an implicit “NTA” judgment: Gloria expresses guilt about lying to her daughter, and Rogers facilitates exploration of these feelings by challenging her self-perception.	17
5.1	Cross-tabulation of Genuineness and Indirect language for the <i>r/aitah</i> dataset. Notice the contrasts in the presence and absence of traits between LLM-generated and human-authored responses.	25
5.2	Cross-tabulation of UPR and Emotional validation for <i>r/aitah</i> dataset. The figure illustrates the stark contrasts in traits of LLM-generated and human-authored responses.	26
5.3	Cross-tabulations for the <i>r/anxiety</i> dataset. Subfigure (b) illustrates the contrasts in community preferences amongst human-authored responses.	27
5.4	Cross-tabulation of UPR and Indirect language for responses to <i>r/anxiety</i> posts. The figure illustrates the consistent and divergent behavior of different language models.	29
5.5	Cross-tabulation of Genuineness and Indirect language for <i>r/anxiety</i> . The figure illustrates the consistent and divergent traits of different language models.	29
5.6	<i>r/aitah</i> —Overall performance scores (Accuracy, Precision, Recall, and F1 Score) for OP-preferred responses and model outputs (GPT-4o, Claude, LLaMA, GPT-oss), comparing model baseline response with CoT dialog insertion. Values represent aggregate scores across the four classes, not per-class results.	30

List of Tables

3.1	Traits associated with the Rogerian PCT and Goffman ToF paradigms, with the latter aligned to [6] to ensure comparability are given below. The prompt definitions of traits for all three paradigms are given in Appendix A.1.	13
3.2	Language Models and Temperature Settings	14
5.1	Agreement scores for each trait. Traits are grouped by evaluative paradigm: ToF and PCT. Accuracy and Cohen’s κ reflect moderate-to-substantial agreement for the majority of traits. Annotations were done by one human familiar with the project and GPT-4o ratings. Note that Moral endorsement (ToF) had poor inter-annotator agreement.	23
5.2	r/anxiety LLM response evaluation on Empathy (PCT) and Sycophancy (ToF) metrics	24
5.3	r/anxiety human performance on select empathy (PCT) and sycophancy (ToF) metrics	24
5.4	r/aitah LLM performance on select Empathy (PCT) and Sycophancy (ToF) traits	28
5.5	r/aitah human performance on select Empathy (PCT) and Sycophancy (ToF) traits	28
5.6	Claude performance on social judgments with and without CoT dialog insertion. . . .	30

Chapter 1

Introduction

Recent advancements in transformer-based large language models (LLMs) have made them increasingly human-like in affective conversation. As virtual assistants, LLMs are increasingly being employed to offer users not only information but also to respond to personal questions [31, 19, 1]. The growing use of LLMs in daily life has sparked important questions about their social implications. The news media have highlighted both the advantages [9] and disadvantages [25] of their use.

This rise in usage that has led to investigations into the behavioral traits LLMs exhibit in their responses to personal questions. As more people rely on these systems for personal support, researchers have begun exploring their long-term effects—particularly along social, psychological, and ethical dimensions [14, 26, 10].

Human responses to another’s personal question are often guided by personal experience and/or behavioral values—be courageous, be helpful, be empathic, and polite. Human exchange and deliberation reflect a broad range of socially shared norms and expectations. There may not be a clear consensus on which value was exhibited in such an interaction. What may seem “empathic” to one person may be a matter of subjective interpretation. As personal queries are increasingly directed to LLMs, we ask: what kinds of behavioral values or traits do these models reflect—and possibly reinforce—through their responses?

It has been noted in NLP research that LLMs and AI models tend to exhibit sycophantic traits [?]. That is, they often agree with users regardless of the accuracy or appropriateness of the statements made. This tendency has been observed and investigated in various contexts [11, 29, 20]. However, “sycophancy” represents just one way of describing the emergent traits of LLMs. Just as humans’ traits may be subjected to interpretational differences, so may be LLMs’. Other studies have also found that LLMs exhibit traits perceived to be more empathic than humans, responding in ways that seem supportive and understanding [16, 26].

Illustrative example of LLM and human responses to a personal query from r/aitah - where AITAH is shorthand for “Am I The Asshole”

Personal Query (Reddit post¹): AITAH for being very mad? I am really mad that my friend is buying me gifts I did not want or need. Even after I told them repeatedly how I feel about this holiday and that I hate gifts?

LLM Response (LLM GPT-4o): Hey there, I totally get why you’re upset. It’s important to communicate your boundaries, and it seems like you’ve done that. Maybe a heart-to-heart conversation could help, try explaining why this is so important to you. Also, try to understand where she’s coming from—some people express love through gift-giving.

Human Response (Reddit Comment²): She should listen to you. You deserve to be listened to. But more concerning is your reaction—it’s not normal to be so emotional about someone not listening to you over what is a relatively harmless thing.

What may be perceived as ‘empathic’ under one paradigm could be critiqued as an instance of ‘social sycophancy’ under another. In the absence of clear normative answers, the same response may be annotated as offering emotional ‘validation’ or as showcasing empathic ‘understanding’. At best it is only a matter of subjective interpretation, however at worst strategies to address LLM social sycophancy may erase constructive and desirable traits of LLM empathy.

Therefore, to examine the behavioral traits of LLMs under these different paradigms in the generation of responses to personal questions and compare them with human authored ones, we present DeepReflect - a system designed to analyze and generate responses through the lens of two distinct psychosocial paradigms.

The system evaluates four language model-generated responses and compares them to two human-authored ones for the same personal query. The personal queries are sourced from the Reddit Archives dataset (r/aitah, r/anxiety). A post forms the body of the personal query for which LLM responses are generated. The two human-authored responses under consideration are: (1) the most upvoted comment (top comment), and (2) the comment the original author (OP) of the post interacted with the most (OP-preferred).

The system is designed to annotate the responses with a binary value to indicate whether a value (i.e., socially exhibited trait) is present or absent. The behavioral values used for annotation are grouped according to three distinct psychosocial paradigms:

1. Rokeach Value Survey (RVS),
2. Carl Rogers’ Person-Centered Therapy (PCT), and
3. Goffman’s Theory of Face (ToF).

¹Paraphrased excerpt of post body from the dataset.

²Paraphrased excerpts of human-authored response from the dataset.

Annotations are created blind to the origin of the response. The project broadens the focus on comparing traits exhibited by LLMs with those observed in human-authored responses under the three distinct paradigms.

Unlike Cheng et al. [6], who focus solely on LLM sycophancy in a social context, this work examines how the selection of different psychosocial paradigms and framing influences the perception of an LLM generated response. It investigates whether the same feature of language models leads to them being perceived as more empathic than humans or whether they are exclusively more sycophantic. Further, it provides discussion on the implications of the divergent connotations associated with the annotating paradigm. The experimental design aims to provide measurements and definitive statistical insight into the values exhibited by the two distinct sources of response in personal conversation.

1.1 Aims and Objectives

The project has dual objectives: (1) to design a technical framework for analyzing and comparing human and LLM responses to personal queries. The system must accommodate multiple psychosocial paradigms, enabling cross-paradigm comparisons; and (2) to conduct a series of experiments investigating the traits of human-written and LLM-generated responses, and evaluating their differences and commonalities.

1.1.1 Research Questions

The context of inputs can substantially shape LLM outputs, influencing not only generations made to address personal questions posed by consumers but also analytical evaluations conducted by researchers. The effect is particularly noteworthy within the LLM-as-a-judge paradigm. As research increasingly highlights patterns and concerns regarding the impacts of LLMs in personal queries and social deliberation, there is a critical need for a framework that can analyze and compare responses across multiple value-based perspectives in social contexts without clear normative answers. The framework also needs to remain extensible for researchers to incorporate additional paradigms as the field and needs evolve. This motivates the following research questions:

RQ1: How can a technical framework that systematically evaluates and compares responses from humans against those generated by LLMs across various psychosocial paradigms be designed?

RQ2: What inter-paradigm insights can the framework yield across different psychosocial frameworks? **RQ2a:** Importantly, to what extent can identical features be annotated with divergent connotations across paradigms—"empathic" under Rogerian PCT versus "sycophantic" under Goffman's ToF?

RQ3: How do LLM-generated responses compare to human-authored responses in the context of personal questions, especially where there are not definitive normative answers?

Finally, we examine how the results may come to influence consumer behavior and broader societal outcomes. We explore a potential control mechanism with chain-of-thought (CoT) reasoning. In this mechanism, we extract a dialog between an expert human and his patient to experimentally determine how such dialogue insertion may affect the LLM response.

Our work enables a systematic comparative analysis of potential benefits and risks, and presents a framework for analysis which can be used by researchers and consumers for leveraging the insights in the intentional design of response LLM generation.

1.2 Key Contributions

The key contributions of this work are: (1) the design and implementation of an extensible framework for analyzing and comparing responses to personal queries across three distinct psychosocial paradigms; (2) a comparative analysis under Rogerian PCT, Goffman's ToF and the RVS framework, illustrating how the choice of paradigm can shape the perception of a response; and (3) insights into the relative strengths and weaknesses of LLM versus human responses, and how these insights can inform the generation of customized responses to personal queries.

1.3 Overview of the Report

The remainder of this report is organized as follows:

Chapter 2 provides an overview of **prior literature**—some of the studies introduced in the introduction—ending with a discussion on closing the gap in the divergent perceptions of LLM responses in affective conversation and their evaluation. This chapter also establishes some of the theoretical foundations for the chosen psychosocial paradigms as well as provides the methodological precedent for evaluating responses by operationalizing psychosocial paradigms with NLP.

Chapter 3 presents the detailed **system design and dataset** underlying the DeepReflect technical framework. It also provides an overview of the operationalization of "empathy" behavioral values and "sycophancy" traits under the two distinct psychosocial paradigms. The design of the system is in response to **RQ1** and enables the experiments and analyses necessary to investigate the subsequent RQs. It also showcases the design of the chain-of-thought (CoT) generation subsystem as a control mechanism. This chapter provides the necessary detail for full replication of the system.

Chapter 4 describes the **experimental setup and methods**, including data preprocessing, and the filtering of comments and posts to create a pairing of human- and LLM-generated responses to personal queries which for the purposes of this project are in the form of Reddit posts. It outlines the CoT generations pipeline, the design of comparative experiments on exhibited values and verdicts (for the r/aitah subreddit), and the use of statistical tests to assess inter- and intra-paradigm differences. The chapter also addresses the validity of the experimental construct.

Chapter 5 presents the **results and analyses** of the experiments, reporting the empirical findings according to the annotation criteria. The results determine cross-paradigm associations, with particular attention to cases where the associations in values of paradigms PCT and ToF produce the results needed to address **RQ2** and **RQ2a** and plots a few of the contingency tables to address **RQ3**. Additionally,

the efficacy of expert informed CoT prompting as a lightweight intervention for steering AI behavior toward more desirable outcomes is analyzed.

Chapter **6 concludes** the report with a **discussion** of the broader implications of the findings on the design of AI assistants and "safety mechanisms". It discusses the epistemic limits of interpreting and evaluating responses in this manner and acknowledges other **limitations**, and directions for future research ending with a synthesis of the contributions and findings of the project.

Chapter 2

Literature Survey

The project sets out on the task to establish a reliable evaluation pipeline with the aim of achieving statistically significant insight into the two divergent interpretations of LLM generated responses in affective conversation. The two interpretations with divergent connotations under consideration are LLM “Sycophancy” and “Empathy” assigned to responses generated in social contexts.

The behaviors “Empathy” and “Sycophancy” may be operationalized for such evaluation with a list of traits associated with them. When a response exhibits that trait, it may be annotated as “present” and “absent” when it does not.

The project sets out on the task to establish a reliable evaluation pipeline with the aim of achieving statistically significant insight into the two divergent interpretations of LLM generated responses in affective conversation. The two interpretations with divergent connotations under consideration are LLM “Sycophancy” and “Empathy” assigned to responses generated in social contexts.

The behaviors “Empathy” and “Sycophancy” may be operationalized for such evaluation with a list of traits associated with them. When a response exhibits that trait, it may be annotated as “present” and “absent” when it does not.

The literature review is broken into three subsections:

1. The theoretical foundations of the Psychosocial paradigms under consideration for the evaluation of a text based response.
2. NLP research relevant to:
 - a. The use of transformer-based LLM chatbots in psychosocial contexts.
 - b. The use of transformer-based LLMs for the design of scalable technical systems, specifically for sample annotations at scale. The project offers the opportunity to investigate the annotations with the LLM-as-judge paradigm.
 - c. Control mechanisms for the phenomenon as part of the response generation.
3. Gaps in the existing literature and how the project aims to cover them.

2.1 Psychosocial paradigms: Theoretical Foundations

2.1.1 Rogerian Person Centered Therapy (PCT)

Carl Rogers outlined six conditions that support therapeutic change, which are commonly summarized in practice as empathy, accurate understanding, unconditional positive regard, and genuineness (also referred to as congruence) (Rogers, 1957). Three of these six qualities are useful for annotating social dialogue because they do not depend on a clinical diagnosis, are not tied to specific content, and can often be identified through linguistic patterns such as expressions of understanding, and a nonjudgmental attitude.

The identification of empathic communication can be operationalized through three core relational conditions that serve as reliable markers of a therapist’s empathic stance. These traits are—genuineness, unconditional positive regard, and accurate understanding [22]:

1. Trait: Genuineness. “The third condition is that the therapist should be, within the confines of this relationship, a congruent, genuine, integrated person. It means that within the relationship he is freely and deeply himself, with his actual experience accurately represented by his awareness of himself. It is the opposite of presenting a facade, either knowingly or unknowingly.”
2. Trait: Unconditional positive regard (UPR). “To the extent that the therapist finds himself experiencing a warm acceptance of each aspect of the client’s experience as being a part of that client, he is experiencing unconditional positive regard.”
3. Trait: Accurate understanding. “The fifth condition is that the therapist is experiencing an accurate, empathic understanding of the client’s awareness of his own experience.”

Based on these three values within the PCT paradigm, the evaluation pipeline can be designed with a list of traits that can be annotated as either present or absent (see Appendix A). This approach helps maintain consistency between different raters and is simple enough to be used in large-scale evaluations.

Cognitive Behavioral Therapy (CBT) was also considered as a potential framework for operationalization. However, CBT relies on structured interventions, goal-oriented dialogue, and content-specific strategies that require contextual nuance for interpretation. These characteristics make it less suitable for the surface-level task of annotating responses [5, p. 69].

A discussion on LLM empathy and its impact on well-being

Research indicates that LLM chatbots are efficacious in establishing the perception of empathy and trust with humans in interaction [10]. The paper shows that while voice-based chatbots were associated with slightly better outcomes than text at average use, greater daily use was linked to lower “socialization with people,” with the neutral-voice condition showing a significant decline at higher usage ($\beta = -0.05$, $p = 0.003$). The findings support designing systems that actively route users to real sources of support (e.g., tooling for connections with real humans like friends/family, empathic handoffs to clinicians or legal support depending on the situation) and that make human care visible in

the product, so users feel the tool was created with their well-being in mind rather than as a substitute for socialization. However, evaluating whether such design choices sustain social connection over time is beyond the scope of this project. The focus of the RQs in this project is to establish a reliable text-based response evaluation pipeline for the two psychosocial paradigms under consideration. Such studies indicate the growing importance of understanding the psychosocial implications of LLM behavior in social contexts.

2.1.2 Goffman’s theory of face (ToF)

Following Goffman’s account of face—the positive self-image people seek to maintain in interaction [12]—we treat LLM social sycophancy as the excessive preservation of a human participant’s face. This operationalization of social sycophancy with Goffman’s theory of face preservation lets us identify when models affirm a user’s self-image or avoid challenging it, even when guidance or correction would be more helpful.

Five face-preserving behaviors (used for the trait annotation):

- **Emotional validation** — reassuring and empathic language that offers comfort without critique.
- **Moral endorsement** — affirming or praising the user’s actions or stance, even when community norms would judge them problematic.
- **Indirect language** — hedging and deference that avoids clear guidance or disagreement.
- **Indirect action** — suggesting coping or reflective steps that do not change the situation when concrete action may be needed.
- **Accepting framing** — taking the user’s assumptions at face value instead of questioning harmful or mistaken premises.

To keep annotations comparable with Cheng et al.’s findings, the operationalization definitions for the five face-preserving behaviors is kept consistent with the ones used in the ELEPHANT framework with some edits to ensure compatibility with the extensible DeepReflect architecture [6].

2.1.3 Rokeach Value Survey (RVS)

Rokeach distinguishes *terminal* (end-states) and *instrumental* (modes of conduct) values and introduces a rank-ordering instrument that has been used to analyze value salience in individuals and groups [24]. For the purposes of this project, the Rokeach’s Value Survey (RVS) lists serve as an additional *interpretive lens* to test the ease of extensibility of the system without limiting the design to the two paradigms under consideration for traits of Empathic and Sycophantic communication.

RVS is divided into two sets of eighteen traits each. The first of which corresponds to traits associated with “modes of conduct” and the second with “terminal (end of life) goals”.

The project is unable to treat the thirty-six RVS value tags with manual annotation validation owing to the large number of traits; instead, this instrument is used to test and design the system for extensibility.

Anthropic’s recent “Values in the Wild” project analyzes hundreds of thousands of real assistant–user conversations to characterize value expressions in deployed systems, offering a useful contemporary value tagging mechanism [13]. However, as a for-profit lab’s internal dataset and coding scheme, the framework has not yet been independently validated by social-scientific experts, and only a subset of artifacts is public; accordingly, we treat it as a non-standard framework only for discussion in this section. A practical limitation for this project is that DeepReflect does not have access to Anthropic’s full proprietary conversation corpus, so comparisons are necessarily approximate and restricted to public releases. This reality further motivates the choice of the RVS paradigm (well-established in social psychology) for the trait annotations.

2.2 Contemporary NLP research

LLM “Empathy” LLM empathy refers to a model’s ability to recognize a user’s emotional state and respond with supportive, context-appropriate language. Both News coverage and academic research on this topic is expanding rapidly [9, 15, 7], with a systematic review cataloging growing evidence that LLMs can express empathy cues and improve user satisfaction in health and support settings [26].

Popular reporting echoes this momentum, noting that companies are actively pursuing “empathetic AI” for customer service and care [2]. In clinical communication, journalists have described clinicians using chatbots to refine bedside manner and draft more compassionate messages (Kolata, 2023). HCI scholars also caution against mistaking convincing displays of emotion for genuine understanding, urging careful design and transparency about system limits (Cuadra et al., 2024). Taken together, these perspectives suggest that LLM empathy—used deliberately and with safeguards—can scale empathic communication to more people and contexts.

LLM “Sycophancy” Sycophancy refers to an LLM chatbot’s tendency to mirror a user’s stated beliefs or preferences—agreeing with or amplifying them even when doing so degrades factual accuracy or contradicts majority social judgment. The phenomenon has gained attention in the social context with the increasing research investigations ([20]). Such studies on LLM sycophancy show that models may agree with plainly incorrect statements once the user signals agreement with no correction. In short, alignment with the preference for affirmations in text based communication can inadvertently bias models toward accommodation over prioritizing truthful information. The issue is further complicated in social reasoning tasks, where models struggle to align with human consensus on socially acceptable behavior, a challenge highlighted in the SocialGaze framework for multi-perspective deliberation [29]. These findings suggest that LLM sycophancy may not only distort factual accuracy but also social norms, reinforcing stereotypes or harmful biases. More recent work systematically evaluates the phenomenon across domains, showing that sycophantic behavior in LLMs can persist in both mathematics and medical advice settings degrading performance in the event of regressive sycophancy [11].

The simultaneous yet independent endeavors in research into the phenomena labeled as LLM sycophancy and LLM empathy motivates the design of an evaluation pipeline that probes both within the same technical framework.

2.2.1 LLM-as-judge paradigm

To design the system for scale, trait annotation is done with the LLM-as-Judge paradigm—with tested prompt definitions. Benchmarking shows that strong judges (e.g., GPT-4o - which is used for this project) can approximate human preferences on multi-turn dialogue (MT-Bench; Chatbot Arena), yet known biases (position, verbosity, self-enhancement) and calibration gaps remain [32, 28]. Accordingly, DeepReflect cross-checks a subset with a human rater to estimate accuracy and interannotator reliability in the social context. This follows broader evaluation practice emphasizing the investigation and benchmarking of annotation errors with this paradigm [11].

Chain-of-thought (CoT) reasoning Finally, the literature review takes note of a design decision in the response generation pipeline: rather than rely on costly fine-tuning to mitigate social failure modes, the system injects *human-expert dialog as a few-shot example* from Carl Roger’s transcribed session with his client, Gloria, that embody the empathic communication skills central to PCT [21]. In the dialog, Rogers exhibits unconditional positive regard without offering an explicit verdict or social judgment to the client.

Chain-of-thought prompting has been shown to elicit multi-step reasoning in sufficiently large models without task-specific fine-tuning, yielding large gains across arithmetic, commonsense, and symbolic tasks (e.g., with PaLM-540B, just eight CoT examples reached state-of-the-art on GSM8K) [30]. With the advantage being that the response generation can be improved by giving a few step-by-step examples. This supports the choice of injecting human expert dialog as a control mechanism for the phenomenon labeled as LLM “Sycophancy” or LLM “Empathy”.

This control mechanism on the four general purpose models under consider (GPToss, GPT 4o, Claude claude-3-5-sonnet-20241022 and Llama meta-llama/Meta-Llama-3.1-70B-Instruct-Turbo) also serves as a preliminary validation for the recent call to focus on the development of expertise informed language models [8].

2.3 Gaps in the Literature and Open Challenges

In sum, as LLM-chatbots have become increasingly human-like and more users seek companionship with them, studies have highlighted both the advantages and disadvantages of their use. While some have raised concerns around “emotional dependence” [10], several others have explored empathic perceptions of LLM responses and found them advantageous not only in the field of medical support and mental health but also in everyday personal queries [16, 26]. However, different psychosocial paradigms tend to frame LLM responses in markedly divergent terms. **What may be perceived as ‘empathy’ under a psychotherapeutic paradigm could instead be critiqued as an instance of ‘social sycophancy’** by frameworks informed by Goffman’s Theory of Face [6]. Importantly, in the absence

of clear normative answers, the same statement may be categorised as ‘face-preserving behaviour’ or ‘unconditional positive regard’.

DeepReflect provides a comparative framework to address this gap by assessing how evaluative judgments are shaped by the psychosocial paradigm through which a response is framed. Moreover, it is designed to be extensible by researchers, enabling the incorporation of both conventional paradigms, such as Rokeach’s values framework, and contemporary discovery-based approaches, such as Anthropic’s Values in the Wild [13], whereas prior work has tended to focus on a single paradigm in isolation.

Finally, the investigation of controlling output response avoids replicating prior work that seeks to mitigate sycophancy exclusively [6]. Instead of treating sycophancy as a defect to be eliminated in isolation, DeepReflect provides a system to situate response generation within extensible psychosocial frameworks. This ensures that outputs are not merely reactive to user prompts but can be guided by well-established instruments for values and personal-growth.

In practice, this involves chain-of-thought reasoning [30] that explicitly incorporates the chosen framework. Unlike approaches that mimic deliberation across hypothetical perspectives [29], this control strategy extends the contractualist, rule-based tradition of questioning developed in [14]. Its key distinction lies in embedding the questioning within expert-informed guidelines. While these prior investigations emphasized plurality of viewpoints and normative exception-handling, this work foregrounds the role of pre-existing psychosocial instruments in shaping the ongoing, ever-changing conversations of personal reflection.

Chapter 3

DeepReflect: System Design

3.1 Dataset

Two datasets were constructed for this project using the Pushshift Reddit Archives [4], originally collected between 2006 and 2023 through the Pushshift API¹. Posts and comments were extracted from two subreddits: (1) r/aitah and (2) r/anxiety. For each post, three components were considered: the body of the original post written by the author (OP), the most upvoted human-written comment (denoted hc1 in Figure 3.1), and the comment with which the OP engaged the most (hc2). Additional detail regarding data filtering and text preprocessing is provided in Section 4. Because the dataset predates the public release of GPT-3.5 in November 2022—and given that large language models (LLMs) only entered widespread public use after early 2023 [17]—all posts and comments in our data are likely majority human-authored.²

3.1.1 Subreddit Selection

The r/anxiety subreddit is a community dedicated to individuals experiencing anxiety and related mental health challenges. Membership does not require a formal diagnosis or medical documentation, which enables broad analyses from psychosocial perspectives. Posts often center on personal struggles, coping strategies and the impact on daily life.

The r/aitah subreddit (short for “Am I The Asshole”) is a community where users seek judgment on personal dilemmas and social interactions. It has over three million members and covers a wide range of topics, including relationships, family dynamics, workplace conflicts, and personal questions. Users typically describe their situations in detail and ask the community to determine whether they were in the wrong (the “asshole”) or not. The crowd-sourced social judgments captured in these posts makes r/aitah a valuable source for examining behaviors and values expressed in digital discussions of personal matters. The crowdsourced verdicts serve as a **proxy for the ground-truth** judgment of the scenario by humans. This is especially valuable for comparing human responses to the situation

¹<https://github.com/pushshift/api>

²Note: Ethical approval for this work was granted by the University Research Ethics Committee of the University of Sheffield (Application Number 068330). See Appendix A.2.1.

against the language model responses under the Goffman’s ToF and Rogerian PCT paradigms which serve as signals for "Sycophantic" and "Empathic" behaviors respectively.

A dataset of 1000 posts was constructed by selecting the top 1000 posts from the *r/aitah* subreddit. Two hundred of these were evenly split between the verdicts: “You’re The Asshole” (YTA) and “Not The Asshole” (NTA), No Verdict and Rest for the CoT experiments. The data was collected from the Pushshift Reddit Archives.

Demographic information at the subreddit level is not available. However, research indicates that Reddit users overall are predominantly American (49.9%), male (67%), and young (22% aged 18–29 years; 14% aged 30–49 years) [3, 27]. While this dataset is not representative of the general population, it reflects a demographic more likely to engage with LLMs for personal queries. This demographic is broadly aligned with the WEIRD (Western, Educated, Industrialized, Rich, Democratic) population, and it must therefore be acknowledged that the results of this study are necessarily constrained to this population.

3.2 DeepReflect

3.2.1 System Design

The system architecture is modular, consisting of two subsystems: (1) the Evaluation Pipeline and (2) the Response Generation Pipeline. A high-level overview is presented in Figure 3.1.

Subsystem 1 is designed to address RQ1 and to be used by researchers interested in the comparative analysis of LLM responses to personal queries across multiple psychosocial paradigms. Four psychosocial paradigms have been implemented in this work. However, the system is designed to be extensible, allowing researchers to incorporate additional paradigms as the field and interests evolve by adding the new paradigm and its associated list of values or behaviors to the system architecture which is then read in during the annotation step.

Subsystem 2 is designed to generate responses to personal queries through a custom-designed chain-of-thought (CoT) reasoning mechanism and can be used by both researchers for analyses (see Section 4) and by consumers for response generation.

Table 3.1: Traits associated with the Rogerian PCT and Goffman ToF paradigms, with the latter aligned to [6] to ensure comparability are given below. The prompt definitions of traits for all three paradigms are given in Appendix A.1.

Paradigm	Traits List
Rogerian PCT (Empathy)	Unconditional positive regard, Genuineness, Accurate Understanding
Goffman ToF (Sycophancy)	Emotional validation, Moral endorsement, Indirect language, Indirect action, Accept framing

Table 3.2: Language Models and Temperature Settings

Language Model	Temperature
GPT-4o ¹	0.7
Claude-3.5 ²	0.7
Llama-2-70B-Chat ³	0.7
GPT-oss ⁴	0.7

Evaluation Framework

The evaluation framework consists of the following steps in a pipeline architecture (see Figure 3.1):




Figure 3.1: Pipeline architecture for DeepReflect.

- 1. Post and Comment extraction:** The top 1000 posts for two subreddits: (1) r/aitah and (2) r/anxiety are extracted from the Reddit Archives dataset. For each post, three components are considered: i. the body of the original post written by the author (OP), ii. the most upvoted human-written comment, and iii. the comment with which the OP engaged the most. Additional detail regarding the top post filtering and text preprocessing are provided in Section 4.
- 2. Basic Language Model Response Generation:** For each post body, a baseline response is generated using an API call to the LLM. This response is appended to a dataframe (p in Figure 3.1) containing: (i) The original post title and body (ii) the top most-upvoted human

¹<https://platform.openai.com/docs/models/gpt-4o>

²<https://www.anthropic.com/index/claude-3-5>

³<https://api.together.xyz/meta-llama/Meta-Llama-3.1-70B-Instruct-Turbo>

⁴[openai/gpt-oss-120b](https://openai.com/gpt-oss-120b)

comment, and (iii) the comment the OP engaged the most with (available for 50% of the posts). The resulting dataset therefore consists of the original post body, paired with two sources of responses to personal queries - human-written and AI responses.

3. **Importing Paradigms and the Associated Values:** The following psychosocial paradigms are implemented in this work: (1) Rogerian PCT, (2) Goffman’s ToF and (3) RVS. Each paradigm is associated with a unique list of values or traits as described in Section 2. The selected paradigms and the associated traits are read into the system for annotations in the next step.
4. **Feature Extraction and Annotation:** Features are extracted and annotated at the full response level for each post and its associated responses, preserving the complete semantic context of the exchange. The annotations are made by GPT-4o with the LLM-as-a-judge [32] paradigm for the three psychosocial paradigms. So, if a response exhibits a value or behavior, it is annotated as **1**, otherwise it is annotated as **0** for each value under the paradigm. For example, features demonstrating “unconditional positive regard,” a value within Rogerian PCT, are annotated as **1** for that value; all others are annotated as **0**.

For the annotation step, human validation is performed with one expert annotator familiar with the research problem. The human annotator annotates on 100 post-response pairs. This validation along with LLM annotations is used to calculate Cohen’s Kappa and accuracy metrics in order to gauge the reliability of the annotations.

$$\kappa = \frac{p_o - p_e}{1 - p_e},$$

p_o = observed agreement (accuracy)

p_e = expected agreement by chance

See section 4 for validation metrics.

5. **Save dataframe to file:** The resulting annotated data, along with the post and corresponding set of responses are saved to a file.
6. **Statistical Analysis:** The annotated dataframe serves as the foundation for subsequent analyses, including (i) comparing value distributions in Reddit versus language model responses across the different paradigms, (ii) conducting topical analyses, and (iii) addressing RQs 2–3 1.1.1 with inter-paradigm correlations.

Note that the standard softmax distribution over a vocabulary of size V for transformer based LLMs with a temperature parameter $T > 0$ that rescales the logits before normalization is:

$$p_i^{(T)} = \frac{e^{z_i/T}}{\sum_{j=1}^V e^{z_j/T}}. \quad (3.1)$$

Lower T ($T < 1$) sharpens the distribution, making the model more deterministic, while higher T ($T > 1$) flattens it, encouraging diversity in the generated responses. For response generations, T is set to $T = 0.7$ to see how responses vary with stochasticity under realistic consumer usage conditions.

DeepReflect Generation Pipeline

In this subsystem, responses to the post are generated through a custom-designed chain-of-thought (CoT) reasoning mechanism. Instead of relying on standard language model outputs, the framework generates responses that are explicitly guided by reasoning chains derived from **expert human reasoning in dialog** and transcripts. The expert human transcripts are retrieved from existing literature within Carl Rogers’ PCT paradigm [21] in this instance. See figure 3.2 for details.

Chain-of-Thought Reasoning

The CoT generation process is formalized as follows:

$$p_{\theta}(y \mid x) = \sum_z p_{\theta}(y \mid x, z) p_{\theta}(z \mid x)$$

where x is the Reddit-based personal query (i.e. a post body), z is the reasoning chain derived from expert human dialog, y is the response generated by DeepReflect and θ denotes the parameters of the base language model. Here, $p_{\theta}(z \mid x)$ denotes the probability distribution over reasoning chains given the query, while $p_{\theta}(y \mid x, z)$ denotes the probability of generating a response conditioned on both the query and reasoning trajectory.

Conditioning on z separates reasoning from surface realization, allowing responses to be shaped by expert-informed CoT patterns rather than unconstrained next-token prediction. See Figure 3.2.

Generated outputs can either be passed through the Evaluation Pipeline for analysis or returned directly in response to a user query. In the case of the former, the project investigates whether PCT-informed CoT reasoning alters verdict distributions (e.g., NTA $\rightarrow$ YTA, No verdict or Rest) and in the case of the latter, a preliminary analysis of response quality is performed given the small sample size (200).

As in the previous section, for evaluation purposes, T is set to both 0 and 1.0 for the CoT generations as well (see Equation 3.1).




Figure 3.2: CoT Generation with personal queries embedded in reasoning dialogs retrieved from expert human transcripts. In this case, the dialogs are from Carl Rogers' sessions with Gloria (patient) [21]. This dialog was selected because it reflects an implicit "NTA" judgment: Gloria expresses guilt about lying to her daughter, and Rogers facilitates exploration of these feelings by challenging her self-perception.

Chapter 4

Methods

4.1 Data preprocessing

A dataset was built from the RedditArchives for two public subreddits—AITAH, and Anxiety. For each subreddit, the top 1,000 most upvoted posts were selected, excluding weekly megathreads, deleted/removed items, and AutoModerator entries. Exact and near-duplicate texts (specifically, crossposts, cypastes and bot repeats) are also removed to prevent inflated counts and biased comparisons. Of these 1000, $N = 897$ posts for r/aitah and $N = 935$ for r/anxiety were found to be usable for the analyses owing to constraints imposed by the API endpoints for the language models.

For every retained post the pipeline extracts (i) the most upvoted comment and (ii) the comment that the OP engaged with the most; all generated and extracted artifacts were saved as CSVs for downstream analysis.

Text was cleaned with minimal, semantics-preserving preprocessing: we removed non-English items, de-identified obvious personal identifiers (usernames and emails), standardized whitespace and Unicode characters, and lightly constrained length (posts 50–500 words; comments 5–300 words) for comparability with language model outputs.

Each pair of post and its associated set of responses is treated as a single analytical unit ('feature'). For the purposes of verdict analysis with this dataset, (discussed in Section ??), the dataset was stratified on verdicts (NTA, YTA, No_Verdict, REST) to ensure balanced representation across the four classes.

The same approach for 'features' are used for the Anxiety dataset for consistency and to ensure comparability. We also treat each feature as a single analytical unit during manual checks, and statistical aggregation.

4.2 Procedures

For each selected post, the system prompts the target language model firstly, with the basic prompt template to generate a response¹. The response is used to establish a **baseline open-ended response**

¹Prompts for this step are provided in the appendices A.2.1

to the post. This response is appended to a table containing: (i) the model-generated response, (ii) the top community upvoted human comment, and (iii) the OP-preferred comment i.e. the comment the OP engaged with (available for approximately 70% of the posts). The resulting dataframe consists of the original post body, paired with two types of responses to personal queries - human and AI responses.

Feature Extraction

- In steps 3 and 4 of the Evaluation Framework (Figure 3.2.1), features² are annotated at the full response level within each body–response pair. For the statistical analysis, these annotations are then aggregated to construct contingency tables, which form the basis of chi-square tests of independence.
- Note that each feature can be annotated with:
 - **Traits exhibited** by the author.

RQ2 focuses on drawing inter- and intra-paradigm comparative insights across the four psychosocial paradigms while sub-research question RQ2a addresses the epistemic limits of interpreting LLM behavior through psychosocial theories in isolation. Specifically, the same feature may be perceived as 'sycophantic' under Goffman's ToF, 'empathic' under Rogerian PCT.

To support these inquiries, the file saved by the evaluation pipeline in step 5 consists of: the annotated features of the original post, annotated features within the set of the 2 different types of responses (human response, language models) for values exhibited or incentivized under the relevant four psychosocial paradigm(s).

This annotated dataset forms the basis for the subsequent analyses necessary to also address RQ3, which studies the differences in the distributions of values between human-authored and language model-generated responses to personal queries.

4.3 Experiments

The experimental design spans two major dimensions: (i) statistical analysis of the annotated response-level features in the form of chi-square tests (ii) class-wise metrics for social verdict categories ('NTA', 'YTA', 'No verdict', 'Rest') across different source types, including one form of human-authored response (top upvoted comment) and four language model responses.

While i. is conducted for each of the two datasets (r/aitah and r/anxiety), under the different psychosocial paradigms, ii. is valid only for the r/aitah dataset, where the responses may or may not contain explicit verdicts - forming four distinct categories (NTA, YTA, No Verdict, Rest).

4.3.1 Experiment 1

In **Experiment 1**, the primary objective is to compare the selected paradigms and analyze the distributions of values across them, with the aim of ultimately determining how paradigm choice can lead to divergent interpretations of the same LLM response.

²Note that in this context, 'feature' refers to the part of the text used for annotation from the post and responses.

While values incentivized are also provided in the results, the analyses are focused on **values exhibited** under each paradigm by the two distinct types of authors of response.

Statistical Methods

The annotated dataset is used to construct contingency tables that shows how two categorical variables co-occur (with the values of a selected paradigm 1 represented across the columns and the values of the second paradigm represented across the rows). Chi-square tests are performed to assess independence between intra- and inter-paradigm values. The Benferroni correction is applied to control the family-wise error rate.

4.3.2 Experiment 2

Experiment 2 is designed to analyze the differences in judgments for the *r/aitah* dataset across two different sources of responses: i. human-authored, and ii. LLM-authored responses for the two psychosocial paradigms - Rogerian PCT and Goffman's ToF.

4.3.3 Statistical Methods

Judgments are extracted from the annotated dataset under four class labels: (i) NTA, (ii) YTA, (iii) No (explicit) verdict and (iv) Rest. These labels are used to construct a 3×3 confusion matrix, with the human-authored judgment as the ground truth and the LLM-authored judgment as the prediction. Per-Class performance metrics and pairwise error rates, including the False Negative Rate (FNR) and False Positive Rate (FPR), are reported for each class label in Section ??.

The measurements thus obtained are used to inform the analysis on how 'judgments' differ between human- and LLM-authored responses.

4.4 CoT Generation Evaluation

Experiments 1 and 2 are rerun, this time to investigate the efficacy of control mechanisms to align the reasoning in language model outputs more closely with that of human experts. The output response of the language model can then be run through the evaluation pipeline again to compare the distributions of values with the human-sourced responses.

The implementation of generations with the chain-of-thought reasoning is detailed under the **Generation Pipeline** subsystem described in Section ??.

4.5 Construct Validity and Evaluation Metrics

To assess construct validity, one human annotator labeled one hundred randomly sampled post-response pairs across all the paradigms for each response type. The PCT paradigm encompasses three traits to operationalize Empathy and Goffman's ToF consists of five used to operationalize Sycophancy. RVS consists of thirty-six values, however for the purposes of this project, this paradigm is included as a proof of concept to demonstrate the extensibility of the technical framework in response to RQ1.

Inter-rater reliability and accuracy for all ToF and PCT metrics with the human annotator and language model (GPT-4o) are reported in Section 5. For the AITAH dataset, verdicts and accompanying statements in responses were used as proxies for the ground truth social verdict.

Chapter 5

Results and Analysis

This section discusses the associations between traits of empathy and sycophancy. With annotations of responses from the Rogerian PCT paradigm used as indicators of Empathy in the response and annotations from Goffman’s ToF used as indicators of sycophancy, the subsequent sections show the results of chi-square test of independence¹ for pairwise combinations of inter-paradigm traits.

For the chi-square tests, H_0 - **the null hypothesis** - is that the traits of the response *are independent* of the joint distribution of responses annotated as trait 1 of Empathy (derived from PCT) and trait 2 of Sycophancy (derived from ToF) i.e. whether a response possesses that trait of empathy tells you nothing about whether it also has the second trait of sycophancy.

And H_1 - **the alternative hypothesis** - is that the two traits under consideration are *not independent* of the joint distribution of responses annotated as trait 1 of Empathy (derived from PCT) and trait 2 of Sycophancy (derived from ToF) i.e., Empathy and Sycophancy trait annotations co-occur systematically (same responses are often annotated both). The level of significance is set at $\alpha = 0.05$ which is standard for tests of independence.

As the annotations are scaled with the LLM-as-judge paradigm, the reliability of the annotations is discussed first before proceeding to the main results of the experiments. The main results of the experiments discuss the level of significance (p-value)² in the association of two traits belonging to different paradigms with divergent connotations and a measure of the strength of that association (ϕ).

5.1 LLM-as-judge Annotation Reliability

Interannotator agreement was calculated using accuracy and Cohen’s κ to assess the reliability of annotations made by GPT-4o against human annotations on a subset of responses for the task under the LLM-as-judge paradigm. Accuracy and inter-annotator agreement fall within the moderate to substantial range for four of five ToF metrics and all three PCT metrics. However, scores for Moral endorsement are notably poor, indicating that pairwise statistical analyses of associations with this

¹The chi-square test of independence is used to test the independence of two categorical - in this case, binary valued - variables.

²A $p < 0.05$ indicates statistically significant association between the two traits leading the results to reject the null hypothesis (H_0).

trait are unreliable. For this reason, in this chapter the primary focus is on the remaining traits of the two paradigms, though the anomalous contingency tables and metrics with Moral endorsement are included in the appendices for reference.

Table 5.1 summarizes the agreement scores and inter-annotator reliability coefficient (Cohen’s κ) for each trait.

Paradigm	Trait	accuracy	κ
Empathy traits (PCT)	UPR (unconditional positive regard)	0.87	0.71
	Genuineness	0.84	0.68
	Accurate Understanding	0.86	0.69
Sycophancy traits (ToF)	Emotional validation	0.88	0.64
	Moral endorsement	0.30	-0.40
	Indirect language	0.79	0.62
	Indirect action	0.80	0.79
	Accept framing	0.75	0.71

Table 5.1: **Agreement scores for each trait.** Traits are grouped by evaluative paradigm: ToF and PCT. Accuracy and Cohen’s κ reflect moderate-to-substantial agreement for the majority of traits. Annotations were done by one human familiar with the project and GPT-4o ratings. Note that Moral endorsement (ToF) had poor inter-annotator agreement.

5.2 "Sycophancy" or "Empathy" ?

This section discusses the pairwise association and independence of metrics under the Rogerian PCT paradigm and the Goffman ToF paradigm³⁴.

5.2.1 Dual Coding of responses as "Empathic" and "Sycophantic"

The chi-square analyses indicate that a single LLM response may be categorized as both "empathic" and socially "sycophantic," depending on the psychosocial paradigm used to annotate it, across all pairwise combinations of Empathic PCT traits and Sycophantic ToF traits for LLM authored responses to posts from *r/anxiety*, and all but one for human authored responses. For example, UPR is a key trait within the PCT paradigm, serving as a marker of empathy. A chi-square test of independence for this variable with emotional validation, which is operationalized as a trait of sycophancy within Goffman’s ToF framework, reveals a statistically significant association ($p < .001$). Table 5.2 reports strong associations between UPR and Emotional validation across models (ϕ values exceeding .908 for *r/anxiety* across all models), suggesting that LLM responses validating emotions are strongly aligned with the empathic criterion of UPR.

³ToF Moral endorsement is excluded from consideration for chi square tests owing to poor inter-annotator agreement and accuracy scores.

⁴All pairwise ϕ and p -values are reported in the appendices and were generated using the IBM SPSS software.

	<i>r/anxiety</i>									
	GPT-oss response		GPT-4o response		Llama response		Claude response		Totals	
	Φ	<i>p</i> -value	Φ	<i>p</i> -value	Φ	<i>p</i> -value	Φ	<i>p</i> -value	Φ	<i>p</i> -value
UPR x Emotional validation	0.908	<.001	0.937	<.001	0.967	<.001	0.923	<.001	0.933	<.001
UPR x Indirect language	0.878	<.001	0.694	<.001	0.524	<.001	0.359	<.001	0.535	<.001
Genuineness x Indirect language	0.552	<.001	0.534	<.001	0.438	<.001	0.382	<.001	0.408	<.001

Table 5.2: *r/anxiety* LLM response evaluation on Empathy (PCT) and Sycophancy (ToF) metrics

	<i>r/anxiety</i>					
	Top comment		OP-preferred		Totals	
	Φ	<i>p</i> -value	Φ	<i>p</i> -value	Φ	<i>p</i> -value
UPR x Emotional validation	0.243	<.001	0.369	<.001	0.304	<.001
UPR x Indirect language	0.200	<.001	0.248	<.001	0.221	<.001
Genuineness x Indirect language	0.022	0.500	0.069	0.059	0.043	0.077

Table 5.3: *r/anxiety* human performance on select empathy (PCT) and sycophancy (ToF) metrics

For some models such as GPT-4o the strength of the association is consistently high indicating that the same response produced by the language model is likely to be annotated as both emotionally validating and as demonstrating unconditional positive **near-deterministically**. These findings highlight the central issue outlined in Section 2.3 and offer direct insight into the question posed by RQ2a (Section 1.1.1).

5.2.2 Trends within human authored responses

The results indicate that while human responses do exhibit associations between empathy-related traits (e.g., unconditional positive regard) and sycophancy-related traits (e.g., emotional validation), the strength of these associations is **markedly lower** than in LLM responses. Notably, some associations in human responses are weak or non-significant (see Genuineness x Indirect language in tables 5.3 and ??), suggesting that the overlap between empathic and sycophantic traits is less pronounced in human-authored content.

For *r/anxiety* (Table 5.2), the associations between UPR and Emotional validation is moderate, with stronger effects in OP-preferred comments ($\phi = 0.369$) compared to top comments ($\phi = 0.243$). This contrasts with LLM responses in Table ??, where the strength of associations for the same measure exceeded $\phi = 0.908$ for every model considered. Similarly, trends can be seen for the *r/aitah* dataset. Taken together, these results show that while human responses demonstrate overlap between empathic and sycophantic traits, the strength of association is lower relative to LLM outputs.

Another noteworthy pattern emerges when comparing top comments with OP-preferred comments in the human authored responses. Although the sample size of responses is smaller for OP-preferred




Figure 5.1: Cross-tabulation of Genuineness and Indirect language for the r/aitah dataset. Notice the contrasts in the presence and absence of traits between LLM-generated and human-authored responses.

comments, owing to the fact that OPs only engaged beyond the post body in 70-75% of the cases, the magnitude of ϕ is scaled by sample size and thus comparisons remain valid.

Across both subreddits, associations are stronger in OP-preferred comments. For example, in *r/anxiety* UPR $\times$ Emotional validation increases from $\phi = 0.243$ in top comments to $\phi = 0.369$ in OP-preferred comments. Similarly, in *r/aitah*, UPR $\times$ Emotional validation rises from $\phi = 0.509$ in top comments to $\phi = 0.544$ in OP-preferred comments.

This suggests that the comments selected by original posters for engagement are more likely to exhibit overlapping the empathic and sycophantic behavioral traits than those receiving the highest community-driven endorsement through upvotes.

Overall, the evaluation of reddit comment responses confirms that associations between empathy- and sycophancy-related traits are present in human-authored responses too, however they occur at lower strengths than in LLM responses. The differentiation in ϕ between top upvoted comment and OP-preferred comments adds an additional layer of analytical insight that is directly observed from the contingency tables of annotated responses as shown in Figure 5.1 and Figure 5.2.

Notes on human preferences in social engagement

Across both *r/anxiety* and *r/aitah*, OP-preferred comments (i.e., those the OP engaged with) show higher ϕ associations for UPR $\times$ Emotional validation than top comments. This suggests that OPs were more likely to interact with responses coded as “empathic” or “sycophantic.” It is worth noting, however, that OP-preferred sets are smaller ($N = 716$ vs. $N = 897$ for top comments in *r/aitah* and $N = 752$ vs. $N = 935$ for top comments in *r/anxiety*).

While these patterns are consistent with Cheng et al.’s observation that preference datasets underlie social “sycophancy,” the present analysis measures associations between annotated traits, whereas




Figure 5.2: Cross-tabulation of UPR and Emotional validation for r/aitah dataset. The figure illustrates the stark contrasts in traits of LLM-generated and human-authored responses.


Cheng et al. examine dataset-level reinforcement learning preferences to determine the root cause of sycophancy [6]. The similarity in these findings is therefore suggestive of a shared underlying trend in preference in responses to personal questions, but it may not be deterministically conclusive. A deeper investigation of this potential connection between individual engagement patterns and large-scale reinforcement learning preference datasets lies beyond the scope of the project.

5.2.3 Strength of association varies by model


Closer examination of the symmetric measures highlights variation across models for the same variables under consideration. For example, for *r/anxiety*, GPT-oss, GPT-4o, and Llama responses show ϕ values above .69 for UPR $\times$ Indirect language but Claude responses show a moderate association at 0.359 for the two traits. In *r/aitah* 5.4, GPT-4o records the strongest association (.699), while Claude again shows the lowest value (0.128), indicating consistently measurable differences in the traits exhibited by the two language models.

Note that although comparatively weaker associations are observed for UPR $\times$ Indirect language and Genuineness $\times$ Indirect language, they are still statistically significant at the 0.05 level across all models. For *r/anxiety*, pooled ϕ values are .535 while for *r/aitah* the corresponding totals for the same variables are .319. These findings suggest that the associated traits are potentially input dependent.

Overall, the results demonstrate that LLM responses consistently show a significant association between empathic and sycophantic features, with particularly stronger associations in the traits than human responses. The distributions of the presence or absence of traits across language models are depicted in Figure ?? and Figure ??.



(a) Genuineness and Accept framing for the *r/aitah* dataset. Note that for this subreddit, both the OP-preferred and community upvoted comments show only a moderate association between the two traits.



(b) UPR and Emotional validation for the *r/aitah* dataset.

Figure 5.3: Cross-tabulations for the *r/anxiety* dataset. Subfigure (b) illustrates the contrasts in community preferences amongst human-authored responses.

5.3 Social Verdicts in *r/aitah*

Since the verdict in the top comment from *r/aitah* can be used as a proxy for the ground truth social judgment of the situation presented by the user, the results showcase accuracy and precision scores comparing LLM judgments to human judgments. For the purposes of this project, the social verdicts were simplified into four categories: YTA (You are the asshole), NTA (Not the asshole), REST (Everyone

	r/aitah									
	GPT-oss response		GPT-4o response		Llama response		Claude response		Totals	
	Φ	<i>p</i> -value	Φ	<i>p</i> -value	Φ	<i>p</i> -value	Φ	<i>p</i> -value	Φ	<i>p</i> -value
UPR x Emotional validation	0.707	<.001	0.939	<.001	0.796	<.001	0.445	<.001	0.642	<.001
UPR x Indirect language	0.464	<.001	0.699	<.001	0.265	<.001	0.128	<.001	0.319	<.001
Genuineness x Indirect language	0.558	<.001	0.559	<.001	0.288	<.001	0.210	<.001	0.302	<.001

Table 5.4: r/aitah LLM performance on select Empathy (PCT) and Sycophancy (ToF) traits

	r/aitah					
	Top comment		OP-preferred		Totals	
	Φ	<i>p</i> -value	Φ	<i>p</i> -value	Φ	<i>p</i> -value
UPR x Emotional validation	0.509	<.001	0.544	<.001	0.528	<.001
UPR x Indirect language	0.236	<.001	0.311	<.001	0.258	<.001
Genuineness x Indirect language	0.070	0.036	0.097	0.010	0.076	0.002

Table 5.5: r/aitah human performance on select Empathy (PCT) and Sycophancy (ToF) traits

sucks here or No assholes here), and No verdict (no clear judgment).

5.3.1 Generations with CoT

Shift in LLM "judgments" with CoT

CoT with dialog insertion was evaluated as a control mechanism, with results varying across models. For Claude, accuracy declined from 0.45 to 0.37 and macro-average F1 from 0.41 to 0.30 with CoT insertion (Table 5.6). At the class level, recall for No verdict rose from 0.12 to 0.76, while recall for YTA fell from 0.48 to 0.08. These results show that CoT dialog insertion redistributed performance across classes rather than producing uniform changes. Full model performances are provided in the appendix.

Across models, CoT dialog insertion did not consistently raise or lower overall metrics such as accuracy or F1. Instead, it shifted the balance of social judgments, with some classes gaining in recall or precision and others showing losses. While the results depicted may provide a useful starting point for investigation in this area, further study is necessary to determine how CoT dialog insertion affects model outputs in the context of personal queries.

Aggregated metrics alone risk obscuring class-level effects and qualitative analyses are required to assess how CoT prompts influence model behavior. The evaluation was conducted on small sample sizes ($N = 200$ per condition per model) due to resource constraints from API costs and the analyses on the raw data necessary to ensure a stratified split. While trends are visible, they should be interpreted cautiously. The implications of limited sample size for generalizability are addressed in the limitations section 6.1.




Figure 5.4: Cross-tabulation of UPR and Indirect language for responses to *r/anxiety* posts. The figure illustrates the consistent and divergent behavior of different language models.




Figure 5.5: Cross-tabulation of Genuineness and Indirect language for *r/anxiety*. The figure illustrates the consistent and divergent traits of different language models.


Qualitative Analysis with CoT

5.3.2 Error analysis

5.3.3 Extensibility of the framework

Social Verdict	With CoT dialog insertion				Without CoT dialog insertion			
	Precision	Recall	F1-score	Support	Precision	Recall	F1-score	Support
No verdict	0.30	0.76	0.43	50	0.60	0.12	0.21	50
NTA	0.48	0.62	0.55	50	0.38	0.96	0.55	50
REST	0.50	0.04	0.07	50	0.32	0.23	0.27	50
YTA	0.67	0.08	0.14	50	0.92	0.48	0.63	50
Accuracy	0.37 (200 comparisons)				0.45 (200 comparisons)			
Macro avg	0.49	0.38	0.30	—	0.55	0.45	0.41	—
Weighted avg	0.49	0.37	0.29	—	0.55	0.45	0.41	—

Table 5.6: Claude performance on social judgments with and without CoT dialog insertion.

Figure 5.6: *r/aitah*—Overall performance scores (Accuracy, Precision, Recall, and F1 Score) for OP-preferred responses and model outputs (GPT-4o, Claude, LLaMA, GPT-oss), comparing model baseline response with CoT dialog insertion. Values represent aggregate scores across the four classes, not per-class results.

Chapter 6

Conclusion

6.1 Limitations

In the course of establishing the reliability of the evaluation pipeline, the following limitations have come to light in order of decreasing precedence:

Firstly, the annotation of features with the LLM-as-judge paradigm remains limited in reliability in cases where the definitions of annotations are semantically nuanced. The issue may not be feasible to fix by increasing context window size as the inaccurate annotation of human authored responses was especially evident in cases of sarcasm (show example). Thus in order to evaluate and annotate human-authored responses may require the consideration of the advancements in sarcasm detection (cite).

In the same theme, the prompt design required multiple iterations and tests to capture the traits to an acceptable degree of reliability – a process which remains susceptible to differences between language models and non-transparent and changing system prompts of the judge model (GPT 4o).

This was particularly limiting in establishing the reliable annotation of the Sycophantic trait - moral endorsement. The definition of moral endorsement was based on the prompt by cheng et al to maintain comparability. While the prompt appeared to work accurately in some cases, when used to annotate the full datasets, it yielded errors. The error analyses indicated that this is likely due to confusion in the semantic interpretation of the definition of the trait, with some batches being annotated such that it captured the opposite trend and some other batches where the trend captured was accurate.

- Moral endorsement: Determine whether the response suggests, implies, or indicates that the question asker (the “I” in the post) did something wrong. • Focus on content, not tone: A response given negatively, sarcastically, or with profanity does not automatically mean the asker did something wrong. REMEMBER: Your goal is to decide only whether the response conveys that the asker’s action, decision, or behavior was wrong (limited strictly to the “I” in the post), regardless of tone.

This issue has led the project to be inconclusive in the pairwise associations with this specific trait. Though it must be noted that the finding is useful in itself, and a key result that contributes to highlighting the limitations of the LLM-as-judge paradigm for such annotations. The use of the paradigm without the step of interannotator reliability can propagate errors in analyses and stand in the way of drawing meaningful conclusions.

Secondly, there may be biases inherent in the datasets formed of the Reddit Archives and within the subreddit communities owing to the majority user demographics and their preferences. Therefore, the human-authored responses used for comparison and to establish a proxy for the ground truth are necessarily subject to the biases inherited from the dataset. OP-preferred comments were only available for about 70% of the posts. Future study may involve the curation of such datasets for research to mitigate these issues.

Lastly, Empathy and Sycophancy are contextually advantageous or disadvantageous in social interactions. While the study looked at Reddit communities which deal primarily with queries on personal matters broadly, Topic analysis of the posts could be especially insightful in identifying specific contexts in which the phenomenon of LLM empathy or sycophancy may be acceptable and even desirable. Rather than a direction of future study, topic analysis on the dataset may be considered a “low hanging fruit” with the current state of findings being limited only by time in this endeavor. The differences in strengths of associations in traits between the two Reddit communities indicate significant potential for topic analyses to further cement the main arguments presented in the following section.

6.2 Summary and Conclusion

This project has demonstrated with evidence that the same phenomenon that researchers and users associate with LLM social “sycophancy” leads to the technology being received and adopted with great ease, making it describable also as an instance of “empathy”.

Further, the research demonstrates that the core issue of divergent connotations being used to describe the trait that points to the same response is not isolated to LLM-authored responses. In social contexts without clear normative answers, the same human authored responses when annotated systematically, may be perceived as sycophantic or empathic too. Perhaps there has not been a profound need to operationalize and measure the difference between sycophancy and empathy in a social context before.

As the use of LLM assistance becomes socially pervasive, owing to capacity of "situationally aware reward hacking (SARH [18])" of human interactions by AI systems, the need to ask such questions of critique and inquiry also grows. What are the measurable traits of empathic and sycophantic communication irrespective of source ? Where lies the nuanced but necessary distinction ?

The findings show that LLM responses exhibit traits that are capable of qualifying them as significantly more empathic than people, a desirable attribute in some interactions. Such highly empathic communication can not only attract and retain new adopters but also be used in specialized fields like education where early accolades and positivity might guide a learner to be more receptive to enduring the tougher paths of complex information. Another area to benefit from this feature could be mental wellbeing, where empathic communication is key. However, even experts in the field of mental wellbeing have prioritized the traits favoring the exchange of truthful information over others that comprise empathy [23, 33]. Empathic LLM communication for its own sake, should it replace various forms of social communication, can become a disservice to the exchange of information. Whether human empathic communication necessarily carries forth this exchange of information is a question

beyond the scope of this current project. Thus, a future line of inquiry may involve establishing the amount of information in human-authored responses and comparing this to the information carried by model generated responses as a measure of the quality of AI generated communication.

As has been established by the field and verified in the course of this project, the same phenomenon of LLM generated responses is categorized as instances of LLM sycophancy. Thus, any lines drawn in the form of “guardrails” or system prompts to mitigate LLM sycophancy could benefit from cautious consideration of the positive effects of the phenomenon. To have none in place, can leave the encounters of end users with LLMs at excessive flattery and spectacle. However, a crude design of mitigation strategies in the form of standard system prompts around LLM sycophancy without the nuanced understanding of the underlying layers from which the phenomenon emerges prohibits access to the benefits of the technology in the form of widely distributable and socially blind empathic communication.

To this end, the work laid down a critical foundation for advancing initial prompt-based control mechanisms with chain-of-thought reasoning based on excerpt sourced transcripts and expert human reasoning dialogs. With analyses on a sample of generated responses indicating qualitative improvements in the generated response. However, any control mechanism of this nature designed to navigate LLM sycophancy or empathy must ensure that it does not stray users too far from that which is beneficial-the exchange of information.

Bibliography

- [1] ANTHROPIC. How people use Claude for support, advice, and companionship, June 2025. Accessed: 2025-08-25.
- [2] BANNON, L. Can AI do empathy even better than humans? companies are trying it. *The Wall Street Journal* (Nov 2023). Accessed on 2023-11-01.
- [3] BARTHEL, M., STOCKING, G., HOLCOMB, J., AND MITCHELL, A. Reddit news users more likely to be male, young and digital in their news preferences. Pew Research Center Report, Feb. 2016.
- [4] BAUMGARTNER, J., ZANNETTOU, S., KEEGAN, B., SQUIRE, M., AND BLACKBURN, J. The Pushshift Reddit dataset. *arXiv preprint arXiv:2001.08435* (2020).
- [5] BECK, J. S. *Cognitive Behavior Therapy: Basics and Beyond*, 2 ed. The Guilford Press, New York, NY, 2011. Foreword by Aaron T. Beck.
- [6] CHENG, M., YU, S., LEE, C., KHADPE, P., IBRAHIM, L., AND JURAFSKY, D. Social sycophancy: A broader understanding of llm sycophancy. *arXiv preprint arXiv:2505.13995* (2025).
- [7] CUADRA, A., WANG, M., STEIN, L. A., JUNG, M. F., DELL, N., ESTRIN, D., AND LANDAY, J. A. The illusion of empathy? Notes on displays of emotion in human-computer interaction. In *Proceedings of the CHI Conference on Human Factors in Computing Systems (CHI '24)* (Honolulu, HI, USA, 2024), ACM, pp. 446:1–446:18.
- [8] DREDZE, M., WINATA, G. I., KAMBADUR, P., WU, S., IRSOY, O., LU, S., DABRAVOLSKI, V., ROSENBERG, D. S., AND GEHRMANN, S. Academics can contribute to domain-specialized language models. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)* (2024), Association for Computational Linguistics, pp. 5100–5110.
- [9] EDITORIAL, G. I asked ChatGPT to listen to me talk about myself – and it actually helped, 2025. The Guardian, Aug 16.
- [10] FANG, C. M., LIU, A. R., VALDEMAR, D., LEE, E., CHAN, S. W. T., PATARANUTAPORN, P., AND MAES, P. How ai and human behaviors shape psychosocial effects of chatbot use: A longitudinal randomized controlled study. *arXiv preprint arXiv:2503.17473 1*, 1 (2025).

- [11] FANOUS, A., GOLDBERG, J., AGARWAL, A. A., LIN, J., ZHOU, A., DANESHJOU, R., AND KOYEJO, S. Syceval: Evaluating llm sycophancy. *arXiv* (Feb. 2025). AIES 2025; v3 last revised 21 Aug 2025.
- [12] GOFFMAN, E. On face-work: An analysis of ritual elements in social interaction. *Psychiatry* 18, 3 (1955), 213–231.
- [13] HUANG, DURMUS, M., ET AL. Values in the wild: Discovering and analyzing values in real-world language model interactions. *arXiv preprint arXiv:2401.00095* (2024).
- [14] JIN, Z., LEVINE, S., GONZALEZ, F. A., KAMAL, O., SAP, M., SACHAN, M., MIHALCEA, R., TENENBAUM, J. B., AND SCHÖLKOPF, B. When to make exceptions: Exploring language models as accounts of human moral judgment. In *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)* (2022). NeurIPS 2022 conference paper; OpenReview version available at OpenReview.
- [15] KOLATA, G. When doctors use a chatbot to improve their bedside manner. *The New York Times* (2023). 2023 article.
- [16] LEE, C., FANG, Y., ZHANG, Y., LIU, Y., WANG, X., LI, X., AND ZHANG, J. Empathic responses in llms: A study of user perceptions. *arXiv preprint arXiv:2505.13995* 1, 1 (2024).
- [17] LIANG, W., ZHANG, Y., CODREANU, M., WANG, J., CAO, H., AND ZOU, J. The widespread adoption of large language model-assisted writing across society. *arXiv preprint arXiv:2502.09747* (2025).
- [18] NGO, R., CHAN, L., AND MINDERMAN, S. The alignment problem from a deep learning perspective. *arXiv preprint* (2022). Revised version as of May 4, 2025; published at ICLR 2024.
- [19] PHANG, J., LAMPE, M., AHMAD, L., AGARWAL, S., FANG, C. M., LIU, A. R., DANRY, V., LEE, E., CHAN, S. W., PATARANUTAPORN, P., AND MAES, P. Investigating affective use and emotional well-being on ChatGPT. Technical report / preprint, OpenAI & MIT Media Lab, Mar. 2025. Accessed: 2025-08-25.
- [20] RANALDI, L., AND PUCCI, G. When large language models contradict humans? large language models’ sycophantic behaviour. *arXiv preprint arXiv:2311.09410* (2025). v4; updated 24 June 2025.
- [21] ROGERS, C. Session with Gloria. In *The Carl Rogers Reader*, H. Kirschenbaum and V. L. Henderson, Eds. Houghton Mifflin, 1989, pp. 198–215. Transcript of Carl Rogers’s counseling session with Gloria, originally filmed in 1965.
- [22] ROGERS, C. R. The necessary and sufficient conditions of therapeutic personality change. *Journal of Consulting Psychology* 21, 2 (1957), 95–103.
- [23] ROGERS, C. R. *On becoming a person: A therapist’s view of psychotherapy*. Houghton Mifflin, Boston, 1961.

- [24] ROKEACH, M. *The Nature of Human Values*. The Free Press, New York, NY, 1973.
- [25] SINGER, N. Teen’s death leads to lawsuit against chatbot company Character.ai, 2024. The New York Times, Oct 23.
- [26] SORIN, V., BRIN, D., BARASH, Y., KONEN, E., CHARNEY, A., NADKARNI, G., AND KLANG, E. Large language models and empathy: Systematic review. *J Med Internet Res* 26 (2024), e52597.
- [27] STATISTA. Reddit global active user distribution. Statista Statistics Portal, 2025. Accessed: 2025-08-24.
- [28] THAKUR, A. S., CHOUDHARY, K., RAMAYAPALLY, V. S., VAIDYANATHAN, S., AND HUPKES, D. Judging the judges: Evaluating alignment and vulnerabilities in LLMs-as-judges. *arXiv* (June 2024).
- [29] VIJJINI, A. R., MENON, R. R., FU, J., SRIVASTAVA, S., AND CHATURVEDI, S. Socialgaze: Improving the integration of human social norms in large language models. *arXiv preprint arXiv:2410.08698*, 2024. Submitted October 11, 2024.
- [30] WEI, J., WANG, X., SCHUURMANS, D., BOSMA, M., ICHTER, B., XIA, F., CHI, E., LE, Q., AND ZHOU, D. Chain-of-thought prompting elicits reasoning in large language models. *arXiv preprint arXiv:2201.11903*, 2022.
- [31] ZHANG, Y., ZHAO, D., HANCOCK, J. T., KRAUT, R., AND YANG, D. The rise of ai companions: How human-chatbot relationships influence well-being. *arXiv preprint arXiv:2506.12605* (June 17 2025). Version 2, submitted on June 14 and revised June 17, 2025.
- [32] ZHENG, L., WANG, Y., LIU, X., ZHANG, H., LI, M., AND ZHANG, J. Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)* (2023), Association for Computational Linguistics.

Appendices

Appendix A

Appendix A

A.1 Annotation Prompts and Traits

A.2 Response Generation Prompts

A.2.1 Basic Response Generation Template



Downloaded: 20/09/2025
Approved: 29/05/2025

Tara Jain
Registration number: 240167262
Computer Science
Programme: Artificial Intelligence MSc - COMT130

Dear Tara

PROJECT TITLE: Digitally Observable Indicators of Declining Mental Health: Trends in anonymized social media posting activity
APPLICATION: Reference Number 068330

On behalf of the University ethics reviewers who reviewed your project, I am pleased to inform you that on 29/05/2025 the above-named project was **approved** on ethics grounds, on the basis that you will adhere to the following documentation that you submitted for ethics review:

- University research ethics application form 068330 (form submission date: 21/05/2025); (expected project end date: 30/09/2025).

If during the course of the project you need to [deviate significantly from the above-approved documentation](#) please inform me since written approval will be required.

Your responsibilities in delivering this research project are set out at the end of this letter.

Yours sincerely

Elizabeth Bridges
Ethics Admin
Computer Science

Please note the following responsibilities of the researcher in delivering the research project:

- The project must abide by the University's Research Ethics Policy: <https://www.sheffield.ac.uk/research-services/ethics-integrity/policy>
- The project must abide by the University's Good Research & Innovation Practices Policy: https://www.sheffield.ac.uk/polopoly_fs/1.671066!/file/GRIPPolicy.pdf
- The researcher must inform their supervisor (in the case of a student) or Ethics Admin (in the case of a member of staff) of any significant changes to the project or the approved documentation.
- The researcher must comply with the requirements of the law and relevant guidelines relating to security and confidentiality of personal data.
- The researcher is responsible for effectively managing the data collected both during and after the end of the project in line with best practice, and any relevant legislative, regulatory or contractual requirements.

Appendix B

Analysis Reports

r/anxiety: Human authored responses

Crosstabulations and Statistical Analysis Results

EMP: UPR * PHANT: Emotional Validation * Comment Filtering

Crosstab					
Count					
Comment Filtering		PHANT: Emotional Validation		Total	
		absent	present		
OP Preferred	EMP: UPR	absent	118	265	383
		present	11	358	369
	Total		129	623	752
Top Upvoted Comment	EMP: UPR	absent	73	393	466
		present	9	460	469
	Total		82	853	935
Total	EMP: UPR	absent	191	658	849
		present	20	818	838
	Total		211	1476	1687

Symmetric Measures				
Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.369	<.001
		Cramer's V	.369	<.001
	N of Valid Cases		752	
Top Upvoted Comment	Nominal by Nominal	Phi	.243	<.001
		Cramer's V	.243	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.304	<.001
		Cramer's V	.304	<.001
	N of Valid Cases		1687	

EMP: UPR * PHANT: Moral Endorsement * Comment Filtering

Crosstab

Count			PHANT: Moral Endorsement		
Comment Filtering			absent	present	Total
OP Preferred	EMP: UPR	absent	355	28	383
		present	368	1	369
	Total		723	29	752
Top Upvoted Comment	EMP: UPR	absent	456	10	466
		present	467	2	469
	Total		923	12	935
Total	EMP: UPR	absent	811	38	849
		present	835	3	838
	Total		1646	41	1687

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	-.183	<.001
		Cramer's V	.183	<.001
	N of Valid Cases		752	
Top Upvoted Comment	Nominal by Nominal	Phi	-.076	.020
		Cramer's V	.076	.020
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	-.134	<.001
		Cramer's V	.134	<.001
	N of Valid Cases		1687	

EMP: UPR * PHANT: Indirect Language * Comment Filtering

Crosstab

Count			PHANT: Indirect Language		
Comment Filtering			absent	present	Total
OP Preferred	EMP: UPR	absent	352	31	383
		present	270	99	369
	Total		622	130	752
Top Upvoted Comment	EMP: UPR	absent	439	27	466
		present	380	89	469
	Total		819	116	935
Total	EMP: UPR	absent	791	58	849
		present	650	188	838
	Total		1441	246	1687

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.248	<.001
		Cramer's V	.248	<.001
	N of Valid Cases		752	
Top Upvoted Comment	Nominal by Nominal	Phi	.200	<.001
		Cramer's V	.200	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.221	<.001
		Cramer's V	.221	<.001
	N of Valid Cases		1687	

EMP: UPR * PHANT: Indirect Action * Comment Filtering

Crosstab

Count			PHANT: Indirect Action		Total
Comment Filtering			absent	present	
OP Preferred	EMP: UPR	absent	231	152	383
		present	209	160	369
	Total		440	312	752
Top Upvoted Comment	EMP: UPR	absent	270	196	466
		present	287	182	469
	Total		557	378	935
Total	EMP: UPR	absent	501	348	849
		present	496	342	838
	Total		997	690	1687

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.037	.307
		Cramer's V	.037	.307
	N of Valid Cases		752	
Top Upvoted Comment	Nominal by Nominal	Phi	-.033	.311
		Cramer's V	.033	.311
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	-.002	.941
		Cramer's V	.002	.941
	N of Valid Cases		1687	

EMP: UPR * PHANT: Accept Framing * Comment Filtering

Crosstab

Count					
			PHANT: Accept Framing		Total
Comment Filtering			absent	present	
OP Preferred	EMP: UPR	absent	60	323	383
		present	9	360	369
	Total		69	683	752
Top Upvoted Comment	EMP: UPR	absent	37	429	466
		present	8	461	469
	Total		45	890	935
Total	EMP: UPR	absent	97	752	849
		present	17	821	838
	Total		114	1573	1687

Symmetric Measures

			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.229	<.001
		Cramer's V	.229	<.001
	N of Valid Cases		752	
Top Upvoted Comment	Nominal by Nominal	Phi	.146	<.001
		Cramer's V	.146	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.187	<.001
		Cramer's V	.187	<.001
	N of Valid Cases		1687	

EMP: Genuineness * PHANT: Emotional Validation * Comment Filtering

Crosstab

Count					
			PHANT: Emotional Validation		Total
Comment Filtering			absent	present	
OP Preferred	EMP: Genuineness	absent	55	12	67
		present	74	611	685
	Total		129	623	752
Top Upvoted Comment	EMP: Genuineness	absent	38	24	62
		present	44	829	873
	Total		82	853	935
Total	EMP: Genuineness	absent	93	36	129
		present	118	1440	1558
	Total		211	1476	1687

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.539	<.001
		Cramer's V	.539	<.001
	N of Valid Cases		752	
Top Upvoted Comment	Nominal by Nominal	Phi	.495	<.001
		Cramer's V	.495	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.518	<.001
		Cramer's V	.518	<.001
	N of Valid Cases		1687	

EMP: Genuineness * PHANT: Moral Endorsement * Comment Filtering

Crosstab

Count			PHANT: Moral Endorsement		Total
Comment Filtering			absent	present	
OP Preferred	EMP: Genuineness	absent	64	3	67
		present	659	26	685
	Total		723	29	752
Top Upvoted Comment	EMP: Genuineness	absent	62	0	62
		present	861	12	873
	Total		923	12	935
Total	EMP: Genuineness	absent	126	3	129
		present	1520	38	1558
	Total		1646	41	1687

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	-.010	.782
		Cramer's V	.010	.782
	N of Valid Cases		752	
Top Upvoted Comment	Nominal by Nominal	Phi	.030	.353
		Cramer's V	.030	.353
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.002	.936
		Cramer's V	.002	.936
	N of Valid Cases		1687	

EMP: Genuineness * PHANT: Indirect Language * Comment Filtering

Crosstab

Count			PHANT: Indirect Language		Total
Comment Filtering			absent	present	
OP Preferred	EMP: Genuineness	absent	61	6	67
		present	561	124	685
	Total		622	130	752
Top Upvoted Comment	EMP: Genuineness	absent	56	6	62
		present	763	110	873
	Total		819	116	935
Total	EMP: Genuineness	absent	117	12	129
		present	1324	234	1558
	Total		1441	246	1687

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.069	.059
		Cramer's V	.069	.059
	N of Valid Cases		752	
Top Upvoted Comment	Nominal by Nominal	Phi	.022	.500
		Cramer's V	.022	.500
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.043	.077
		Cramer's V	.043	.077
	N of Valid Cases		1687	

EMP: Genuineness * PHANT: Indirect Action * Comment Filtering

Crosstab

Count			PHANT: Indirect Action		Total
Comment Filtering			absent	present	
OP Preferred	EMP: Genuineness	absent	50	17	67
		present	390	295	685
	Total		440	312	752
Top Upvoted Comment	EMP: Genuineness	absent	48	14	62
		present	509	364	873
	Total		557	378	935
Total	EMP: Genuineness	absent	98	31	129
		present	899	659	1558
	Total		997	690	1687

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.102	.005
		Cramer's V	.102	.005
	N of Valid Cases		752	
Top Upvoted Comment	Nominal by Nominal	Phi	.097	.003
		Cramer's V	.097	.003
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.099	<.001
		Cramer's V	.099	<.001
	N of Valid Cases		1687	

EMP: Genuineness * PHANT: Accept Framing * Comment Filtering

Crosstab

Count			PHANT: Accept Framing		Total
Comment Filtering		absent	present		
OP Preferred	EMP: Genuineness	absent	39	28	67
		present	30	655	685
	Total		69	683	752
Top Upvoted Comment	EMP: Genuineness	absent	32	30	62
		present	13	860	873
	Total		45	890	935
Total	EMP: Genuineness	absent	71	58	129
		present	43	1515	1558
	Total		114	1573	1687

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.531	<.001
		Cramer's V	.531	<.001
	N of Valid Cases		752	
Top Upvoted Comment	Nominal by Nominal	Phi	.583	<.001
		Cramer's V	.583	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.553	<.001
		Cramer's V	.553	<.001
	N of Valid Cases		1687	

EMP: Understanding * PHANT: Emotional Validation * Comment Filtering

Crosstab

Count			PHANT: Emotional Validation		Total
Comment Filtering			absent	present	
OP Preferred	EMP: Understanding	absent	108	95	203
		present	21	528	549
	Total		129	623	752
Top Upvoted Comment	EMP: Understanding	absent	60	100	160
		present	22	753	775
	Total		82	853	935
Total	EMP: Understanding	absent	168	195	363
		present	43	1281	1324
	Total		211	1476	1687

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.581	<.001
		Cramer's V	.581	<.001
	N of Valid Cases		752	
Top Upvoted Comment	Nominal by Nominal	Phi	.462	<.001
		Cramer's V	.462	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.535	<.001
		Cramer's V	.535	<.001
	N of Valid Cases		1687	

EMP: Understanding * PHANT: Moral Endorsement * Comment Filtering

Crosstab

Count			PHANT: Moral Endorsement		Total
Comment Filtering			absent	present	
OP Preferred	EMP: Understanding	absent	183	20	203
		present	540	9	549
	Total		723	29	752
Top Upvoted Comment	EMP: Understanding	absent	156	4	160
		present	767	8	775
	Total		923	12	935
Total	EMP: Understanding	absent	339	24	363
		present	1307	17	1324
	Total		1646	41	1687

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	-.189	<.001
		Cramer's V	.189	<.001
	N of Valid Cases		752	
Top Upvoted Comment	Nominal by Nominal	Phi	-.049	.133
		Cramer's V	.049	.133
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	-.142	<.001
		Cramer's V	.142	<.001
	N of Valid Cases		1687	

EMP: Understanding * PHANT: Indirect Language * Comment Filtering

Crosstab

Count		PHANT: Indirect Language		Total
Comment Filtering		absent	present	
OP Preferred	EMP: Understanding	absent	184	203
		present	438	549
	Total		622	752
Top Upvoted Comment	EMP: Understanding	absent	154	160
		present	665	775
	Total		819	935
Total	EMP: Understanding	absent	338	363
		present	1103	1324
	Total		1441	1687

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.127	<.001
		Cramer's V	.127	<.001
	N of Valid Cases		752	
Top Upvoted Comment	Nominal by Nominal	Phi	.119	<.001
		Cramer's V	.119	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.114	<.001
		Cramer's V	.114	<.001
	N of Valid Cases		1687	

EMP: Understanding * PHANT: Indirect Action * Comment Filtering

Crosstab

Count					
			PHANT: Indirect Action		Total
Comment Filtering			absent	present	
OP Preferred	EMP: Understanding	absent	147	56	203
		present	293	256	549
	Total		440	312	752
Top Upvoted Comment	EMP: Understanding	absent	125	35	160
		present	432	343	775
	Total		557	378	935
Total	EMP: Understanding	absent	272	91	363
		present	725	599	1324
	Total		997	690	1687

Symmetric Measures

			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.172	<.001
		Cramer's V	.172	<.001
	N of Valid Cases		752	
Top Upvoted Comment	Nominal by Nominal	Phi	.172	<.001
		Cramer's V	.172	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.169	<.001
		Cramer's V	.169	<.001
	N of Valid Cases		1687	

EMP: Understanding * PHANT: Accept Framing * Comment Filtering

Crosstab

Count					
			PHANT: Accept Framing		Total
Comment Filtering			absent	present	
OP Preferred	EMP: Understanding	absent	55	148	203
		present	14	535	549
	Total		69	683	752
Top Upvoted Comment	EMP: Understanding	absent	34	126	160
		present	11	764	775
	Total		45	890	935
Total	EMP: Understanding	absent	89	274	363
		present	25	1299	1324
	Total		114	1573	1687

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.377	<.001
		Cramer's V	.377	<.001
	N of Valid Cases		752	
Top Upvoted Comment	Nominal by Nominal	Phi	.349	<.001
		Cramer's V	.349	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.370	<.001
		Cramer's V	.370	<.001
	N of Valid Cases		1687	

r/aitah: Human authored responses

Crosstabulations and Statistical Analysis Results

EMP: UPR * PHANT: Emotional Validation * Comment Filtering

Crosstab					
Count					
			PHANT: Emotional Validation		
Comment Filtering			absent	present	Total
OP Preferred	EMP: UPR	absent	413	137	550
		present	20	146	166
	Total		433	283	716
Top Upvoted Comment	EMP: UPR	absent	403	209	612
		present	32	253	285
	Total		435	462	897
Total	EMP: UPR	absent	816	346	1162
		present	52	399	451
	Total		868	745	1613

Symmetric Measures				
Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.544	<.001
		Cramer's V	.544	<.001
	N of Valid Cases		716	
Top Upvoted Comment	Nominal by Nominal	Phi	.509	<.001
		Cramer's V	.509	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.528	<.001
		Cramer's V	.528	<.001
	N of Valid Cases		1613	

EMP: UPR * PHANT: Moral Endorsement * Comment Filtering

Crosstab

Count			PHANT: Moral Endorsement		
Comment Filtering			absent	present	Total
OP Preferred	EMP: UPR	absent	236	314	550
		present	144	22	166
	Total		380	336	716
Top Upvoted Comment	EMP: UPR	absent	309	303	612
		present	267	18	285
	Total		576	321	897
Total	EMP: UPR	absent	545	617	1162
		present	411	40	451
	Total		956	657	1613

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	-.371	<.001
		Cramer's V	.371	<.001
	N of Valid Cases		716	
Top Upvoted Comment	Nominal by Nominal	Phi	-.420	<.001
		Cramer's V	.420	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	-.404	<.001
		Cramer's V	.404	<.001
	N of Valid Cases		1613	

EMP: UPR * PHANT: Indirect Language * Comment Filtering

Crosstab

Count			PHANT: Indirect Language		
Comment Filtering			absent	present	Total
OP Preferred	EMP: UPR	absent	489	61	550
		present	101	65	166
	Total		590	126	716
Top Upvoted Comment	EMP: UPR	absent	572	40	612
		present	220	65	285
	Total		792	105	897
Total	EMP: UPR	absent	1061	101	1162
		present	321	130	451
	Total		1382	231	1613

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.311	<.001
		Cramer's V	.311	<.001
	N of Valid Cases		716	
Top Upvoted Comment	Nominal by Nominal	Phi	.236	<.001
		Cramer's V	.236	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.258	<.001
		Cramer's V	.258	<.001
	N of Valid Cases		1613	

EMP: UPR * PHANT: Indirect Action * Comment Filtering

Crosstab

Count			PHANT: Indirect Action		Total
Comment Filtering			absent	present	
OP Preferred	EMP: UPR	absent	433	117	550
		present	97	69	166
	Total		530	186	716
Top Upvoted Comment	EMP: UPR	absent	531	81	612
		present	197	88	285
	Total		728	169	897
Total	EMP: UPR	absent	964	198	1162
		present	294	157	451
	Total		1258	355	1613

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.195	<.001
		Cramer's V	.195	<.001
	N of Valid Cases		716	
Top Upvoted Comment	Nominal by Nominal	Phi	.210	<.001
		Cramer's V	.210	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.193	<.001
		Cramer's V	.193	<.001
	N of Valid Cases		1613	

EMP: UPR * PHANT: Accept Framing * Comment Filtering

Crosstab

Count					
			PHANT: Accept Framing		Total
Comment Filtering			absent	present	
OP Preferred	EMP: UPR	absent	173	377	550
		present	7	159	166
	Total		180	536	716
Top Upvoted Comment	EMP: UPR	absent	124	488	612
		present	6	279	285
	Total		130	767	897
Total	EMP: UPR	absent	297	865	1162
		present	13	438	451
	Total		310	1303	1613

Symmetric Measures

			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.265	<.001
		Cramer's V	.265	<.001
	N of Valid Cases		716	
Top Upvoted Comment	Nominal by Nominal	Phi	.240	<.001
		Cramer's V	.240	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.258	<.001
		Cramer's V	.258	<.001
	N of Valid Cases		1613	

EMP: Genuineness * PHANT: Emotional Validation * Comment Filtering

Crosstab

Count					
			PHANT: Emotional Validation		Total
Comment Filtering			absent	present	
OP Preferred	EMP: Genuineness	absent	63	11	74
		present	370	272	642
	Total		433	283	716
Top Upvoted Comment	EMP: Genuineness	absent	39	8	47
		present	396	454	850
	Total		435	462	897
Total	EMP: Genuineness	absent	102	19	121
		present	766	726	1492
	Total		868	745	1613

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.171	<.001
		Cramer's V	.171	<.001
	N of Valid Cases		716	
Top Upvoted Comment	Nominal by Nominal	Phi	.162	<.001
		Cramer's V	.162	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.174	<.001
		Cramer's V	.174	<.001
	N of Valid Cases		1613	

EMP: Genuineness * PHANT: Moral Endorsement * Comment Filtering

Crosstab

Count			PHANT: Moral Endorsement		Total
Comment Filtering			absent	present	
OP Preferred	EMP: Genuineness	absent	60	14	74
		present	320	322	642
	Total		380	336	716
Top Upvoted Comment	EMP: Genuineness	absent	44	3	47
		present	532	318	850
	Total		576	321	897
Total	EMP: Genuineness	absent	104	17	121
		present	852	640	1492
	Total		956	657	1613

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.191	<.001
		Cramer's V	.191	<.001
	N of Valid Cases		716	
Top Upvoted Comment	Nominal by Nominal	Phi	.144	<.001
		Cramer's V	.144	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.155	<.001
		Cramer's V	.155	<.001
	N of Valid Cases		1613	

EMP: Genuineness * PHANT: Indirect Language * Comment Filtering

Crosstab

Count			PHANT: Indirect Language		
Comment Filtering			absent	present	Total
OP Preferred	EMP: Genuineness	absent	69	5	74
		present	521	121	642
	Total		590	126	716
Top Upvoted Comment	EMP: Genuineness	absent	46	1	47
		present	746	104	850
	Total		792	105	897
Total	EMP: Genuineness	absent	115	6	121
		present	1267	225	1492
	Total		1382	231	1613

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.097	.010
		Cramer's V	.097	.010
	N of Valid Cases		716	
Top Upvoted Comment	Nominal by Nominal	Phi	.070	.036
		Cramer's V	.070	.036
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.076	.002
		Cramer's V	.076	.002
	N of Valid Cases		1613	

EMP: Genuineness * PHANT: Indirect Action * Comment Filtering

Crosstab

Count			PHANT: Indirect Action		
Comment Filtering			absent	present	Total
OP Preferred	EMP: Genuineness	absent	60	14	74
		present	470	172	642
	Total		530	186	716
Top Upvoted Comment	EMP: Genuineness	absent	40	7	47
		present	688	162	850
	Total		728	169	897
Total	EMP: Genuineness	absent	100	21	121
		present	1158	334	1492
	Total		1258	355	1613

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.055	.144
		Cramer's V	.055	.144
	N of Valid Cases		716	
Top Upvoted Comment	Nominal by Nominal	Phi	.024	.477
		Cramer's V	.024	.477
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.032	.199
		Cramer's V	.032	.199
	N of Valid Cases		1613	

EMP: Genuineness * PHANT: Accept Framing * Comment Filtering

Crosstab

Count			PHANT: Accept Framing		Total
Comment Filtering			absent	present	
OP Preferred	EMP: Genuineness	absent	40	34	74
		present	140	502	642
	Total		180	536	716
Top Upvoted Comment	EMP: Genuineness	absent	31	16	47
		present	99	751	850
	Total		130	767	897
Total	EMP: Genuineness	absent	71	50	121
		present	239	1253	1492
	Total		310	1303	1613

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.226	<.001
		Cramer's V	.226	<.001
	N of Valid Cases		716	
Top Upvoted Comment	Nominal by Nominal	Phi	.344	<.001
		Cramer's V	.344	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.285	<.001
		Cramer's V	.285	<.001
	N of Valid Cases		1613	

EMP: Understanding * PHANT: Emotional Validation * Comment Filtering

Crosstab

Count				
		PHANT: Emotional Validation		Total
Comment Filtering		absent	present	
OP Preferred	EMP: Understanding	.00	212	235
		1.00	221	481
	Total		433	716
Top Upvoted Comment	EMP: Understanding	.00	126	140
		1.00	309	757
	Total		435	897
Total	EMP: Understanding	.00	338	375
		1.00	530	1238
	Total		868	1613

Symmetric Measures

			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.425	<.001
		Cramer's V	.425	<.001
	N of Valid Cases		716	
Top Upvoted Comment	Nominal by Nominal	Phi	.357	<.001
		Cramer's V	.357	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.401	<.001
		Cramer's V	.401	<.001
	N of Valid Cases		1613	

EMP: Understanding * PHANT: Moral Endorsement * Comment Filtering

Crosstab

Count				
		PHANT: Moral Endorsement		Total
Comment Filtering		absent	present	
OP Preferred	EMP: Understanding	.00	115	235
		1.00	265	481
	Total		380	716
Top Upvoted Comment	EMP: Understanding	.00	73	140
		1.00	503	757
	Total		576	897
Total	EMP: Understanding	.00	188	375
		1.00	768	1238
	Total		956	1613

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	-.058	.121
		Cramer's V	.058	.121
	N of Valid Cases		716	
Top Upvoted Comment	Nominal by Nominal	Phi	-.108	.001
		Cramer's V	.108	.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	-.102	<.001
		Cramer's V	.102	<.001
	N of Valid Cases		1613	

EMP: Understanding * PHANT: Indirect Language * Comment Filtering

Crosstab

Count			PHANT: Indirect Language		Total
Comment Filtering			absent	present	
OP Preferred	EMP: Understanding	.00	218	17	235
		1.00	372	109	481
	Total		590	126	716
Top Upvoted Comment	EMP: Understanding	.00	138	2	140
		1.00	654	103	757
	Total		792	105	897
Total	EMP: Understanding	.00	356	19	375
		1.00	1026	212	1238
	Total		1382	231	1613

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.190	<.001
		Cramer's V	.190	<.001
	N of Valid Cases		716	
Top Upvoted Comment	Nominal by Nominal	Phi	.137	<.001
		Cramer's V	.137	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.145	<.001
		Cramer's V	.145	<.001
	N of Valid Cases		1613	

EMP: Understanding * PHANT: Indirect Action * Comment Filtering

Crosstab

Count					
			PHANT: Indirect Action		
Comment Filtering			absent	present	Total
OP Preferred	EMP: Understanding	.00	180	55	235
		1.00	350	131	481
	Total		530	186	716
Top Upvoted Comment	EMP: Understanding	.00	114	26	140
		1.00	614	143	757
	Total		728	169	897
Total	EMP: Understanding	.00	294	81	375
		1.00	964	274	1238
	Total		1258	355	1613

Symmetric Measures

			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.041	.272
		Cramer's V	.041	.272
	N of Valid Cases		716	
Top Upvoted Comment	Nominal by Nominal	Phi	.003	.929
		Cramer's V	.003	.929
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.005	.827
		Cramer's V	.005	.827
	N of Valid Cases		1613	

EMP: Understanding * PHANT: Accept Framing * Comment Filtering

Crosstab

Count					
		PHANT: Accept Framing			
Comment Filtering			absent	present	Total
OP Preferred	EMP: Understanding	.00	99	136	235
		1.00	81	400	481
	Total		180	536	716
Top Upvoted Comment	EMP: Understanding	.00	60	80	140
		1.00	70	687	757
	Total		130	767	897
Total	EMP: Understanding	.00	159	216	375
		1.00	151	1087	1238
	Total		310	1303	1613

Symmetric Measures

Comment Filtering			Value	Approximate Significance
OP Preferred	Nominal by Nominal	Phi	.274	<.001
		Cramer's V	.274	<.001
	N of Valid Cases		716	
Top Upvoted Comment	Nominal by Nominal	Phi	.347	<.001
		Cramer's V	.347	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.324	<.001
		Cramer's V	.324	<.001
	N of Valid Cases		1613	

r/ aith: LLM authored responses

Crosstabulations and Statistical Analysis Results

EMP: UPR * PHANT: Emotional Validation * Language Model

Crosstab

Count					
			PHANT: Emotional Validation		Total
Language Model			absent	present	
claude	EMP: UPR	absent	39	115	154
		present	3	740	743
	Total		42	855	897
gpt-4o	EMP: UPR	absent	31	4	35
		present	0	862	862
	Total		31	866	897
gpt-oss	EMP: UPR	absent	30	28	58
		present	0	837	837
	Total		30	865	895
llama	EMP: UPR	absent	31	17	48
		present	0	849	849
	Total		31	866	897
Total	EMP: UPR	absent	131	164	295
		present	3	3288	3291
	Total		134	3452	3586

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.445	<.001
		Cramer's V	.445	<.001
	N of Valid Cases		897	
gpt-4o	Nominal by Nominal	Phi	.939	<.001
		Cramer's V	.939	<.001
	N of Valid Cases		897	
gpt-oss	Nominal by Nominal	Phi	.707	<.001
		Cramer's V	.707	<.001
	N of Valid Cases		895	
llama	Nominal by Nominal	Phi	.796	<.001
		Cramer's V	.796	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.642	<.001
		Cramer's V	.642	<.001
	N of Valid Cases		3586	

EMP: UPR * PHANT: Moral Endorsement * Language Model

Crosstab

Count					
			PHANT: Moral Endorsement		Total
Language Model			absent	present	
claude	EMP: UPR	absent	91	63	154
		present	686	57	743
	Total		777	120	897
gpt-4o	EMP: UPR	absent	34	1	35
		present	815	47	862
	Total		849	48	897
gpt-oss	EMP: UPR	absent	38	20	58
		present	716	121	837
	Total		754	141	895
llama	EMP: UPR	absent	35	13	48
		present	751	98	849
	Total		786	111	897
Total	EMP: UPR	absent	198	97	295
		present	2968	323	3291
	Total		3166	420	3586

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	-.368	<.001
		Cramer's V	.368	<.001
	N of Valid Cases		897	
gpt-4o	Nominal by Nominal	Phi	.022	.504
		Cramer's V	.022	.504
	N of Valid Cases		897	
gpt-oss	Nominal by Nominal	Phi	-.135	<.001
		Cramer's V	.135	<.001
	N of Valid Cases		895	
llama	Nominal by Nominal	Phi	-.106	.001
		Cramer's V	.106	.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	-.197	<.001
		Cramer's V	.197	<.001
	N of Valid Cases		3586	

EMP: UPR * PHANT: Indirect Language * Language Model

Crosstab

Count					
			PHANT: Indirect Language		Total
Language Model			absent	present	
claude	EMP: UPR	absent	92	62	154
		present	318	425	743
	Total		410	487	897
gpt-4o	EMP: UPR	absent	31	4	35
		present	23	839	862
	Total		54	843	897
gpt-oss	EMP: UPR	absent	32	26	58
		present	38	799	837
	Total		70	825	895
llama	EMP: UPR	absent	34	14	48
		present	177	672	849
	Total		211	686	897
Total	EMP: UPR	absent	189	106	295
		present	556	2735	3291
	Total		745	2841	3586

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.128	<.001
		Cramer's V	.128	<.001
	N of Valid Cases		897	
gpt-4o	Nominal by Nominal	Phi	.699	<.001
		Cramer's V	.699	<.001
	N of Valid Cases		897	
gpt-oss	Nominal by Nominal	Phi	.464	<.001
		Cramer's V	.464	<.001
	N of Valid Cases		895	
llama	Nominal by Nominal	Phi	.265	<.001
		Cramer's V	.265	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.319	<.001
		Cramer's V	.319	<.001
	N of Valid Cases		3586	

EMP: UPR * PHANT: Indirect Action * Language Model

Crosstab

Count			PHANT: Indirect Action		
Language Model			absent	present	Total
claude	EMP: UPR	absent	122	32	154
		present	529	214	743
	Total		651	246	897
gpt-4o	EMP: UPR	absent	31	4	35
		present	54	808	862
	Total		85	812	897
gpt-oss	EMP: UPR	absent	44	14	58
		present	374	463	837
	Total		418	477	895
llama	EMP: UPR	absent	40	8	48
		present	276	573	849
	Total		316	581	897
Total	EMP: UPR	absent	237	58	295
		present	1233	2058	3291
	Total		1470	2116	3586

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.068	.042
		Cramer's V	.068	.042
	N of Valid Cases		897	
gpt-4o	Nominal by Nominal	Phi	.544	<.001
		Cramer's V	.544	<.001
	N of Valid Cases		897	
gpt-oss	Nominal by Nominal	Phi	.154	<.001
		Cramer's V	.154	<.001
	N of Valid Cases		895	
llama	Nominal by Nominal	Phi	.239	<.001
		Cramer's V	.239	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.240	<.001
		Cramer's V	.240	<.001
	N of Valid Cases		3586	

EMP: UPR * PHANT: Accept Framing * Language Model

Crosstab

Count					
			PHANT: Accept Framing		Total
Language Model			absent	present	
claude	EMP: UPR	absent	65	89	154
		present	31	712	743
	Total		96	801	897
gpt-4o	EMP: UPR	absent	32	3	35
		present	15	847	862
	Total		47	850	897
gpt-oss	EMP: UPR	absent	40	18	58
		present	24	813	837
	Total		64	831	895
llama	EMP: UPR	absent	38	10	48
		present	24	825	849
	Total		62	835	897
Total	EMP: UPR	absent	175	120	295
		present	94	3197	3291
	Total		269	3317	3586

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.464	<.001
		Cramer's V	.464	<.001
	N of Valid Cases		897	
gpt-4o	Nominal by Nominal	Phi	.779	<.001
		Cramer's V	.779	<.001
	N of Valid Cases		897	
gpt-oss	Nominal by Nominal	Phi	.632	<.001
		Cramer's V	.632	<.001
	N of Valid Cases		895	
llama	Nominal by Nominal	Phi	.677	<.001
		Cramer's V	.677	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.589	<.001
		Cramer's V	.589	<.001
	N of Valid Cases		3586	

EMP: Genuiness * PHANT: Emotional Validation * Language Model

Crosstab

Count			PHANT: Emotional Validation		Total
Language Model			absent	present	
claude	EMP: Genuiness	absent	30	2	32
		present	12	853	865
	Total		42	855	897
gpt-4o	EMP: Genuiness	absent	31	21	52
		present	0	845	845
	Total		31	866	897
gpt-oss	EMP: Genuiness	absent	30	8	38
		present	0	857	857
	Total		30	865	895
llama	EMP: Genuiness	absent	31	11	42
		present	0	855	855
	Total		31	866	897
Total	EMP: Genuiness	absent	122	42	164
		present	12	3410	3422
	Total		134	3452	3586

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.811	<.001
		Cramer's V	.811	<.001
	N of Valid Cases		897	
gpt-4o	Nominal by Nominal	Phi	.763	<.001
		Cramer's V	.763	<.001
	N of Valid Cases		897	
gpt-oss	Nominal by Nominal	Phi	.884	<.001
		Cramer's V	.884	<.001
	N of Valid Cases		895	
llama	Nominal by Nominal	Phi	.854	<.001
		Cramer's V	.854	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.816	<.001
		Cramer's V	.816	<.001
	N of Valid Cases		3586	

EMP: Genuiness * PHANT: Moral Endorsement * Language Model

Crosstab

Count		PHANT: Moral Endorsement			
Language Model			absent	present	Total
claude	EMP: Genuiness	absent	32	0	32
		present	745	120	865
	Total		777	120	897
gpt-4o	EMP: Genuiness	absent	52	0	52
		present	797	48	845
	Total		849	48	897
gpt-oss	EMP: Genuiness	absent	38	0	38
		present	716	141	857
	Total		754	141	895
llama	EMP: Genuiness	absent	37	5	42
		present	749	106	855
	Total		786	111	897
Total	EMP: Genuiness	absent	159	5	164
		present	3007	415	3422
	Total		3166	420	3586

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.076	.024
		Cramer's V	.076	.024
	N of Valid Cases		897	
gpt-4o	Nominal by Nominal	Phi	.059	.077
		Cramer's V	.059	.077
	N of Valid Cases		897	
gpt-oss	Nominal by Nominal	Phi	.091	.006
		Cramer's V	.091	.006
	N of Valid Cases		895	
llama	Nominal by Nominal	Phi	.003	.925
		Cramer's V	.003	.925
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.059	<.001
		Cramer's V	.059	<.001
	N of Valid Cases		3586	

EMP: Genuiness * PHANT: Indirect Language * Language Model

Crosstab

Count			PHANT: Indirect Language		
Language Model			absent	present	Total
claude	EMP: Genuiness	absent	32	0	32
		present	378	487	865
	Total		410	487	897
gpt-4o	EMP: Genuiness	absent	31	21	52
		present	23	822	845
	Total		54	843	897
gpt-oss	EMP: Genuiness	absent	30	8	38
		present	40	817	857
	Total		70	825	895
llama	EMP: Genuiness	absent	33	9	42
		present	178	677	855
	Total		211	686	897
Total	EMP: Genuiness	absent	126	38	164
		present	619	2803	3422
	Total		745	2841	3586

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.210	<.001
		Cramer's V	.210	<.001
	N of Valid Cases		897	
gpt-4o	Nominal by Nominal	Phi	.559	<.001
		Cramer's V	.559	<.001
	N of Valid Cases		897	
gpt-oss	Nominal by Nominal	Phi	.558	<.001
		Cramer's V	.558	<.001
	N of Valid Cases		895	
llama	Nominal by Nominal	Phi	.288	<.001
		Cramer's V	.288	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.302	<.001
		Cramer's V	.302	<.001
	N of Valid Cases		3586	

EMP: Genuiness * PHANT: Indirect Action * Language Model

Crosstab

Count			PHANT: Indirect Action		
Language Model			absent	present	Total
claude	EMP: Genuiness	absent	31	1	32
		present	620	245	865
	Total		651	246	897
gpt-4o	EMP: Genuiness	absent	32	20	52
		present	53	792	845
	Total		85	812	897
gpt-oss	EMP: Genuiness	absent	32	6	38
		present	386	471	857
	Total		418	477	895
llama	EMP: Genuiness	absent	36	6	42
		present	280	575	855
	Total		316	581	897
Total	EMP: Genuiness	absent	131	33	164
		present	1339	2083	3422
	Total		1470	2116	3586

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.105	.002
		Cramer's V	.105	.002
	N of Valid Cases		897	
gpt-4o	Nominal by Nominal	Phi	.441	<.001
		Cramer's V	.441	<.001
	N of Valid Cases		897	
gpt-oss	Nominal by Nominal	Phi	.158	<.001
		Cramer's V	.158	<.001
	N of Valid Cases		895	
llama	Nominal by Nominal	Phi	.234	<.001
		Cramer's V	.234	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.173	<.001
		Cramer's V	.173	<.001
	N of Valid Cases		3586	

EMP: Genuiness * PHANT: Accept Framing * Language Model

Crosstab

Count			PHANT: Accept Framing		Total
Language Model			absent	present	
claude	EMP: Genuiness	absent	30	2	32
		present	66	799	865
	Total		96	801	897
gpt-4o	EMP: Genuiness	absent	32	20	52
		present	15	830	845
	Total		47	850	897
gpt-oss	EMP: Genuiness	absent	30	8	38
		present	34	823	857
	Total		64	831	895
llama	EMP: Genuiness	absent	31	11	42
		present	31	824	855
	Total		62	835	897
Total	EMP: Genuiness	absent	123	41	164
		present	146	3276	3422
	Total		269	3317	3586

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.517	<.001
		Cramer's V	.517	<.001
	N of Valid Cases		897	
gpt-4o	Nominal by Nominal	Phi	.627	<.001
		Cramer's V	.627	<.001
	N of Valid Cases		897	
gpt-oss	Nominal by Nominal	Phi	.587	<.001
		Cramer's V	.587	<.001
	N of Valid Cases		895	
llama	Nominal by Nominal	Phi	.585	<.001
		Cramer's V	.585	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.561	<.001
		Cramer's V	.561	<.001
	N of Valid Cases		3586	

EMP: Understanding * PHANT: Emotional Validation * Language Model

Crosstab

Count					
			PHANT: Emotional Validation		Total
Language Model			absent	present	
claude	EMP: Understanding	absent	32	2	34
		present	10	863	863
	Total		42	865	897
gpt-4o	EMP: Understanding	absent	31	3	34
		present	0	863	863
	Total		31	866	897
gpt-oss	EMP: Understanding	absent	30	2	32
		present	0	863	863
	Total		30	865	895
llama	EMP: Understanding	absent	31	4	35
		present	0	862	862
	Total		31	866	897
Total	EMP: Understanding	absent	124	11	135
		present	10	3441	3451
	Total		134	3452	3586

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.840	<.001
		Cramer's V	.840	<.001
	N of Valid Cases		897	
gpt-4o	Nominal by Nominal	Phi	.953	<.001
		Cramer's V	.953	<.001
	N of Valid Cases		897	
gpt-oss	Nominal by Nominal	Phi	.967	<.001
		Cramer's V	.967	<.001
	N of Valid Cases		895	
llama	Nominal by Nominal	Phi	.939	<.001
		Cramer's V	.939	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.919	<.001
		Cramer's V	.919	<.001
	N of Valid Cases		3586	

EMP: Understanding * PHANT: Moral Endorsement * Language Model

Crosstab

Count					
			PHANT: Moral Endorsement		Total
Language Model			absent	present	
claude	EMP: Understanding	absent	34	0	34
		present	743	120	863
	Total		777	120	897
gpt-4o	EMP: Understanding	absent	34	0	34
		present	815	48	863
	Total		849	48	897
gpt-oss	EMP: Understanding	absent	32	0	32
		present	722	141	863
	Total		754	141	895
llama	EMP: Understanding	absent	32	3	35
		present	754	108	862
	Total		786	111	897
Total	EMP: Understanding	absent	132	3	135
		present	3034	417	3451
	Total		3166	420	3586

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.078	.019
		Cramer's V	.078	.019
	N of Valid Cases		897	
gpt-4o	Nominal by Nominal	Phi	.047	.158
		Cramer's V	.047	.158
	N of Valid Cases		897	
gpt-oss	Nominal by Nominal	Phi	.083	.013
		Cramer's V	.083	.013
	N of Valid Cases		895	
llama	Nominal by Nominal	Phi	.023	.486
		Cramer's V	.023	.486
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.058	<.001
		Cramer's V	.058	<.001
	N of Valid Cases		3586	

EMP: Understanding * PHANT: Indirect Language * Language Model

Crosstab

Count					
			PHANT: Indirect Language		Total
Language Model			absent	present	
claude	EMP: Understanding	absent	32	2	34
		present	378	485	863
	Total		410	487	897
gpt-4o	EMP: Understanding	absent	31	3	34
		present	23	840	863
	Total		54	843	897
gpt-oss	EMP: Understanding	absent	30	2	32
		present	40	823	863
	Total		70	825	895
llama	EMP: Understanding	absent	32	3	35
		present	179	683	862
	Total		211	686	897
Total	EMP: Understanding	absent	125	10	135
		present	620	2831	3451
	Total		745	2841	3586

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.193	<.001
		Cramer's V	.193	<.001
	N of Valid Cases		897	
gpt-4o	Nominal by Nominal	Phi	.711	<.001
		Cramer's V	.711	<.001
	N of Valid Cases		897	
gpt-oss	Nominal by Nominal	Phi	.616	<.001
		Cramer's V	.616	<.001
	N of Valid Cases		895	
llama	Nominal by Nominal	Phi	.323	<.001
		Cramer's V	.323	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.350	<.001
		Cramer's V	.350	<.001
	N of Valid Cases		3586	

EMP: Understanding * PHANT: Indirect Action * Language Model

Crosstab

Count					
			PHANT: Indirect Action		Total
Language Model			absent	present	
claude	EMP: Understanding	absent	31	3	34
		present	620	243	863
	Total		651	246	897
gpt-4o	EMP: Understanding	absent	31	3	34
		present	54	809	863
	Total		85	812	897
gpt-oss	EMP: Understanding	absent	31	1	32
		present	387	476	863
	Total		418	477	895
llama	EMP: Understanding	absent	33	2	35
		present	283	579	862
	Total		316	581	897
Total	EMP: Understanding	absent	126	9	135
		present	1344	2107	3451
	Total		1470	2116	3586

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.083	.013
		Cramer's V	.083	.013
	N of Valid Cases		897	
gpt-4o	Nominal by Nominal	Phi	.554	<.001
		Cramer's V	.554	<.001
	N of Valid Cases		897	
gpt-oss	Nominal by Nominal	Phi	.194	<.001
		Cramer's V	.194	<.001
	N of Valid Cases		895	
llama	Nominal by Nominal	Phi	.249	<.001
		Cramer's V	.249	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.210	<.001
		Cramer's V	.210	<.001
	N of Valid Cases		3586	

EMP: Understanding * PHANT: Accept Framing * Language Model

Crosstab

Count					
			PHANT: Accept Framing		Total
Language Model			absent	present	
claude	EMP: Understanding	absent	30	4	34
		present	66	797	863
	Total		96	801	897
gpt-4o	EMP: Understanding	absent	31	3	34
		present	16	847	863
	Total		47	850	897
gpt-oss	EMP: Understanding	absent	30	2	32
		present	34	829	863
	Total		64	831	895
llama	EMP: Understanding	absent	31	4	35
		present	31	831	862
	Total		62	835	897
Total	EMP: Understanding	absent	122	13	135
		present	147	3304	3451
	Total		269	3317	3586

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.498	<.001
		Cramer's V	.498	<.001
	N of Valid Cases		897	
gpt-4o	Nominal by Nominal	Phi	.766	<.001
		Cramer's V	.766	<.001
	N of Valid Cases		897	
gpt-oss	Nominal by Nominal	Phi	.647	<.001
		Cramer's V	.647	<.001
	N of Valid Cases		895	
llama	Nominal by Nominal	Phi	.649	<.001
		Cramer's V	.649	<.001
	N of Valid Cases		897	
Total	Nominal by Nominal	Phi	.622	<.001
		Cramer's V	.622	<.001
	N of Valid Cases		3586	

r/anxiety: LLM authored Responses

Crosstabulations and Statistical Analysis Results

EMP: UPR * PHANT: Emotional Validation * Language Model

Crosstab

Count					
			PHANT: Emotional Validation		Total
Language Model			absent	present	
claude	EMP: UPR	absent	30	5	35
		present	0	900	900
	Total		30	905	935
gpt-4o	EMP: UPR	absent	30	4	34
		present	0	901	901
	Total		30	905	935
gpt-oss	EMP: UPR	absent	30	5	35
		present	1	898	899
	Total		31	903	934
llama	EMP: UPR	absent	30	2	32
		present	0	903	903
	Total		30	905	935
Total	EMP: UPR	absent	120	16	136
		present	1	3602	3603
	Total		121	3618	3739

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.923	<.001
		Cramer's V	.923	<.001
	N of Valid Cases		935	
gpt-4o	Nominal by Nominal	Phi	.937	<.001
		Cramer's V	.937	<.001
	N of Valid Cases		935	
gpt-oss	Nominal by Nominal	Phi	.908	<.001
		Cramer's V	.908	<.001
	N of Valid Cases		934	
llama	Nominal by Nominal	Phi	.967	<.001
		Cramer's V	.967	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.933	<.001
		Cramer's V	.933	<.001
	N of Valid Cases		3739	

EMP: UPR * PHANT: Moral Endorsement * Language Model

Crosstab

Count			PHANT: Moral Endorsement		
Language Model			absent	present	Total
claude	EMP: UPR	absent	35	0	35
		present	899	1	900
	Total		934	1	935
gpt-4o	EMP: UPR	absent	34		34
		present	901		901
	Total		935		935
gpt-oss	EMP: UPR	absent	35		35
		present	899		899
	Total		934		934
llama	EMP: UPR	absent	32		32
		present	903		903
	Total		935		935
Total	EMP: UPR	absent	136	0	136
		present	3602	1	3603
	Total		3738	1	3739

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.006	.844
		Cramer's V	.006	.844
	N of Valid Cases		935	
gpt-4o	Nominal by Nominal	Phi	. ^c	
	N of Valid Cases		935	
gpt-oss	Nominal by Nominal	Phi	. ^c	
	N of Valid Cases		934	
llama	Nominal by Nominal	Phi	. ^c	
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.003	.846
		Cramer's V	.003	.846
	N of Valid Cases		3739	

c. No statistics are computed because PHANT: Moral Endorsement is a constant.

EMP: UPR * PHANT: Indirect Language * Language Model

Crosstab

Count			PHANT: Indirect Language		
Language Model			absent	present	Total
claude	EMP: UPR	absent	30	5	35
		present	130	770	900
	Total		160	775	935
gpt-4o	EMP: UPR	absent	30	4	34
		present	23	878	901
	Total		53	882	935
gpt-oss	EMP: UPR	absent	30	5	35
		present	3	896	899
	Total		33	901	934
llama	EMP: UPR	absent	30	2	32
		present	64	839	903
	Total		94	841	935
Total	EMP: UPR	absent	120	16	136
		present	220	3383	3603
	Total		340	3399	3739

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.359	<.001
		Cramer's V	.359	<.001
	N of Valid Cases		935	
gpt-4o	Nominal by Nominal	Phi	.694	<.001
		Cramer's V	.694	<.001
	N of Valid Cases		935	
gpt-oss	Nominal by Nominal	Phi	.878	<.001
		Cramer's V	.878	<.001
	N of Valid Cases		934	
llama	Nominal by Nominal	Phi	.524	<.001
		Cramer's V	.524	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.535	<.001
		Cramer's V	.535	<.001
	N of Valid Cases		3739	

EMP: UPR * PHANT: Indirect Action * Language Model

Crosstab

Count			PHANT: Indirect Action		Total
Language Model			absent	present	
claude	EMP: UPR	absent	32	3	35
		present	162	738	900
	Total		194	741	935
gpt-4o	EMP: UPR	absent	30	4	34
		present	21	880	901
	Total		51	884	935
gpt-oss	EMP: UPR	absent	31	4	35
		present	176	723	899
	Total		207	727	934
llama	EMP: UPR	absent	30	2	32
		present	52	851	903
	Total		82	853	935
Total	EMP: UPR	absent	123	13	136
		present	411	3192	3603
	Total		534	3205	3739

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.344	<.001
		Cramer's V	.344	<.001
	N of Valid Cases		935	
gpt-4o	Nominal by Nominal	Phi	.708	<.001
		Cramer's V	.708	<.001
	N of Valid Cases		935	
gpt-oss	Nominal by Nominal	Phi	.315	<.001
		Cramer's V	.315	<.001
	N of Valid Cases		934	
llama	Nominal by Nominal	Phi	.566	<.001
		Cramer's V	.566	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.423	<.001
		Cramer's V	.423	<.001
	N of Valid Cases		3739	

EMP: UPR * PHANT: Accept Framing * Language Model

Crosstab

Count					
			PHANT: Accept Framing		Total
Language Model			absent	present	
claude	EMP: UPR	absent	30	5	35
		present	46	854	900
	Total		76	859	935
gpt-4o	EMP: UPR	absent	30	4	34
		present	15	886	901
	Total		45	890	935
gpt-oss	EMP: UPR	absent	30	5	35
		present	41	858	899
	Total		71	863	934
llama	EMP: UPR	absent	30	2	32
		present	65	838	903
	Total		95	840	935
Total	EMP: UPR	absent	120	16	136
		present	167	3436	3603
	Total		287	3452	3739

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.560	<.001
		Cramer's V	.560	<.001
	N of Valid Cases		935	
gpt-4o	Nominal by Nominal	Phi	.757	<.001
		Cramer's V	.757	<.001
	N of Valid Cases		935	
gpt-oss	Nominal by Nominal	Phi	.582	<.001
		Cramer's V	.582	<.001
	N of Valid Cases		934	
llama	Nominal by Nominal	Phi	.521	<.001
		Cramer's V	.521	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.588	<.001
		Cramer's V	.588	<.001
	N of Valid Cases		3739	

EMP: Genuiness * PHANT: Emotional Validation * Language Model

Crosstab

Count		PHANT: Emotional Validation			
Language Model			absent	present	Total
claude	EMP: Genuiness	absent	30	4	34
		present	0	901	901
	Total		30	905	935
gpt-4o	EMP: Genuiness	absent	30	24	54
		present	0	881	881
	Total		30	905	935
gpt-oss	EMP: Genuiness	absent	30	53	83
		present	1	850	851
	Total		31	903	934
llama	EMP: Genuiness	absent	30	18	48
		present	0	887	887
	Total		30	905	935
Total	EMP: Genuiness	absent	120	99	219
		present	1	3519	3520
	Total		121	3618	3739

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.937	<.001
		Cramer's V	.937	<.001
	N of Valid Cases		935	
gpt-4o	Nominal by Nominal	Phi	.735	<.001
		Cramer's V	.735	<.001
	N of Valid Cases		935	
gpt-oss	Nominal by Nominal	Phi	.572	<.001
		Cramer's V	.572	<.001
	N of Valid Cases		934	
llama	Nominal by Nominal	Phi	.783	<.001
		Cramer's V	.783	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.727	<.001
		Cramer's V	.727	<.001
	N of Valid Cases		3739	

EMP: Genuiness * PHANT: Moral Endorsement * Language Model

Crosstab

Count		PHANT: Moral Endorsement			
Language Model			absent	present	Total
claude	EMP: Genuiness	absent	34	0	34
		present	900	1	901
	Total		934	1	935
gpt-4o	EMP: Genuiness	absent	54		54
		present	881		881
	Total		935		935
gpt-oss	EMP: Genuiness	absent	83		83
		present	851		851
	Total		934		934
llama	EMP: Genuiness	absent	48		48
		present	887		887
	Total		935		935
Total	EMP: Genuiness	absent	219	0	219
		present	3519	1	3520
	Total		3738	1	3739

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.006	.846
		Cramer's V	.006	.846
	N of Valid Cases		935	
gpt-4o	Nominal by Nominal	Phi	. ^c	
	N of Valid Cases		935	
	Nominal by Nominal	Phi	. ^c	
gpt-oss	Nominal by Nominal	Phi	. ^c	
	N of Valid Cases		934	
	Nominal by Nominal	Phi	. ^c	
llama	Nominal by Nominal	Phi	. ^c	
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.004	.803
		Cramer's V	.004	.803
	N of Valid Cases		3739	

c. No statistics are computed because PHANT: Moral Endorsement is a constant.

EMP: Genuiness * PHANT: Indirect Language * Language Model

Crosstab

Count			PHANT: Indirect Language		Total
Language Model			absent	present	
claude	EMP: Genuiness	absent	31	3	34
		present	129	772	901
	Total		160	775	935
gpt-4o	EMP: Genuiness	absent	30	24	54
		present	23	858	881
	Total		53	882	935
gpt-oss	EMP: Genuiness	absent	30	53	83
		present	3	848	851
	Total		33	901	934
llama	EMP: Genuiness	absent	32	16	48
		present	62	825	887
	Total		94	841	935
Total	EMP: Genuiness	absent	123	96	219
		present	217	3303	3520
	Total		340	3399	3739

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.382	<.001
		Cramer's V	.382	<.001
	N of Valid Cases		935	
gpt-4o	Nominal by Nominal	Phi	.534	<.001
		Cramer's V	.534	<.001
	N of Valid Cases		935	
gpt-oss	Nominal by Nominal	Phi	.552	<.001
		Cramer's V	.552	<.001
	N of Valid Cases		934	
llama	Nominal by Nominal	Phi	.438	<.001
		Cramer's V	.438	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.408	<.001
		Cramer's V	.408	<.001
	N of Valid Cases		3739	

EMP: Genuiness * PHANT: Indirect Action * Language Model

Crosstab

Count			PHANT: Indirect Action		
Language Model			absent	present	Total
claude	EMP: Genuiness	absent	32	2	34
		present	162	739	901
	Total		194	741	935
gpt-4o	EMP: Genuiness	absent	30	24	54
		present	21	860	881
	Total		51	884	935
gpt-oss	EMP: Genuiness	absent	37	46	83
		present	170	681	851
	Total		207	727	934
llama	EMP: Genuiness	absent	31	17	48
		present	51	836	887
	Total		82	853	935
Total	EMP: Genuiness	absent	130	89	219
		present	404	3116	3520
	Total		534	3205	3739

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.351	<.001
		Cramer's V	.351	<.001
	N of Valid Cases		935	
gpt-4o	Nominal by Nominal	Phi	.546	<.001
		Cramer's V	.546	<.001
	N of Valid Cases		935	
gpt-oss	Nominal by Nominal	Phi	.169	<.001
		Cramer's V	.169	<.001
	N of Valid Cases		934	
llama	Nominal by Nominal	Phi	.459	<.001
		Cramer's V	.459	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.321	<.001
		Cramer's V	.321	<.001
	N of Valid Cases		3739	

EMP: Genuiness * PHANT: Accept Framing * Language Model

Crosstab

Count					
			PHANT: Accept Framing		Total
Language Model			absent	present	
claude	EMP: Genuiness	absent	30	4	34
		present	46	855	901
	Total		76	859	935
gpt-4o	EMP: Genuiness	absent	30	24	54
		present	15	866	881
	Total		45	890	935
gpt-oss	EMP: Genuiness	absent	31	52	83
		present	40	811	851
	Total		71	863	934
llama	EMP: Genuiness	absent	31	17	48
		present	64	823	887
	Total		95	840	935
Total	EMP: Genuiness	absent	122	97	219
		present	165	3355	3520
	Total		287	3452	3739

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.569	<.001
		Cramer's V	.569	<.001
	N of Valid Cases		935	
gpt-4o	Nominal by Nominal	Phi	.587	<.001
		Cramer's V	.587	<.001
	N of Valid Cases		935	
gpt-oss	Nominal by Nominal	Phi	.351	<.001
		Cramer's V	.351	<.001
	N of Valid Cases		934	
llama	Nominal by Nominal	Phi	.419	<.001
		Cramer's V	.419	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.450	<.001
		Cramer's V	.450	<.001
	N of Valid Cases		3739	

EMP: Understanding * PHANT: Emotional Validation * Language Model

Crosstab

Count					
			PHANT: Emotional Validation		Total
Language Model			absent	present	
claude	EMP: Understanding	absent	30	2	32
		present	0	903	903
	Total		30	905	935
gpt-4o	EMP: Understanding	absent	30	4	34
		present	0	901	901
	Total		30	905	935
gpt-oss	EMP: Understanding	absent	30	4	34
		present	1	899	900
	Total		31	903	934
llama	EMP: Understanding	absent	30	2	32
		present	0	903	903
	Total		30	905	935
Total	EMP: Understanding	absent	120	12	132
		present	1	3606	3607
	Total		121	3618	3739

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.967	<.001
		Cramer's V	.967	<.001
	N of Valid Cases		935	
gpt-4o	Nominal by Nominal	Phi	.937	<.001
		Cramer's V	.937	<.001
	N of Valid Cases		935	
gpt-oss	Nominal by Nominal	Phi	.921	<.001
		Cramer's V	.921	<.001
	N of Valid Cases		934	
llama	Nominal by Nominal	Phi	.967	<.001
		Cramer's V	.967	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.948	<.001
		Cramer's V	.948	<.001
	N of Valid Cases		3739	

EMP: Understanding * PHANT: Moral Endorsement * Language Model

Crosstab

Count					
			PHANT: Moral Endorsement		Total
Language Model			absent	present	
claude	EMP: Understanding	absent	32	0	32
		present	902	1	903
	Total		934	1	935
gpt-4o	EMP: Understanding	absent	34		34
		present	901		901
	Total		935		935
gpt-oss	EMP: Understanding	absent	34		34
		present	900		900
	Total		934		934
llama	EMP: Understanding	absent	32		32
		present	903		903
	Total		935		935
Total	EMP: Understanding	absent	132	0	132
		present	3606	1	3607
	Total		3738	1	3739

Symmetric Measures

Language Model				Value	Approximate Significance
claude	Nominal by Nominal	Phi		.006	.851
		Cramer's V		.006	.851
	N of Valid Cases			935	
gpt-4o	Nominal by Nominal	Phi		. ^c	
	N of Valid Cases			935	
gpt-oss	Nominal by Nominal	Phi		. ^c	
	N of Valid Cases			934	
llama	Nominal by Nominal	Phi		. ^c	
	N of Valid Cases			935	
Total	Nominal by Nominal	Phi		.003	.848
		Cramer's V		.003	.848
	N of Valid Cases			3739	

c. No statistics are computed because PHANT: Moral Endorsement is a constant.

EMP: Understanding * PHANT: Indirect Language * Language Model

Crosstab

Count					
			PHANT: Indirect Language		Total
Language Model			absent	present	
claude	EMP: Understanding	absent	30	2	32
		present	130	773	903
	Total		160	775	935
gpt-4o	EMP: Understanding	absent	30	4	34
		present	23	878	901
	Total		53	882	935
gpt-oss	EMP: Understanding	absent	30	4	34
		present	3	897	900
	Total		33	901	934
llama	EMP: Understanding	absent	30	2	32
		present	64	839	903
	Total		94	841	935
Total	EMP: Understanding	absent	120	12	132
		present	220	3387	3607
	Total		340	3399	3739

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.383	<.001
		Cramer's V	.383	<.001
	N of Valid Cases		935	
gpt-4o	Nominal by Nominal	Phi	.694	<.001
		Cramer's V	.694	<.001
	N of Valid Cases		935	
gpt-oss	Nominal by Nominal	Phi	.892	<.001
		Cramer's V	.892	<.001
	N of Valid Cases		934	
llama	Nominal by Nominal	Phi	.524	<.001
		Cramer's V	.524	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.544	<.001
		Cramer's V	.544	<.001
	N of Valid Cases		3739	

EMP: Understanding * PHANT: Indirect Action * Language Model

Crosstab

Count					
			PHANT: Indirect Action		Total
Language Model			absent	present	
claude	EMP: Understanding	absent	31	1	32
		present	163	740	903
	Total		194	741	935
gpt-4o	EMP: Understanding	absent	30	4	34
		present	21	880	901
	Total		51	884	935
gpt-oss	EMP: Understanding	absent	30	4	34
		present	177	723	900
	Total		207	727	934
llama	EMP: Understanding	absent	30	2	32
		present	52	851	903
	Total		82	853	935
Total	EMP: Understanding	absent	121	11	132
		present	413	3194	3607
	Total		534	3205	3739

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.353	<.001
		Cramer's V	.353	<.001
	N of Valid Cases		935	
gpt-4o	Nominal by Nominal	Phi	.708	<.001
		Cramer's V	.708	<.001
	N of Valid Cases		935	
gpt-oss	Nominal by Nominal	Phi	.309	<.001
		Cramer's V	.309	<.001
	N of Valid Cases		934	
llama	Nominal by Nominal	Phi	.566	<.001
		Cramer's V	.566	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.423	<.001
		Cramer's V	.423	<.001
	N of Valid Cases		3739	

EMP: Understanding * PHANT: Accept Framing * Language Model

Crosstab

Count					
			PHANT: Accept Framing		Total
Language Model			absent	present	
claude	EMP: Understanding	absent	30	2	32
		present	46	857	903
	Total		76	859	935
gpt-4o	EMP: Understanding	absent	30	4	34
		present	15	886	901
	Total		45	890	935
gpt-oss	EMP: Understanding	absent	30	4	34
		present	41	859	900
	Total		71	863	934
llama	EMP: Understanding	absent	30	2	32
		present	65	838	903
	Total		95	840	935
Total	EMP: Understanding	absent	120	12	132
		present	167	3440	3607
	Total		287	3452	3739

Symmetric Measures

Language Model			Value	Approximate Significance
claude	Nominal by Nominal	Phi	.590	<.001
		Cramer's V	.590	<.001
	N of Valid Cases		935	
gpt-4o	Nominal by Nominal	Phi	.757	<.001
		Cramer's V	.757	<.001
	N of Valid Cases		935	
gpt-oss	Nominal by Nominal	Phi	.591	<.001
		Cramer's V	.591	<.001
	N of Valid Cases		934	
llama	Nominal by Nominal	Phi	.521	<.001
		Cramer's V	.521	<.001
	N of Valid Cases		935	
Total	Nominal by Nominal	Phi	.598	<.001
		Cramer's V	.598	<.001
	N of Valid Cases		3739	